



# Red Hat

## Red Hat Enterprise Linux 10

### Deploying and managing RHEL on Google Cloud

Obtaining RHEL system images and creating RHEL instances on Google Cloud



# Red Hat Enterprise Linux 10 Deploying and managing RHEL on Google Cloud

---

Obtaining RHEL system images and creating RHEL instances on Google Cloud

## Legal Notice

Copyright © Red Hat.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat Software Collections is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

## Abstract

To use Red Hat Enterprise Linux (RHEL) on Google Cloud, you can create and deploy RHEL system images. This also involves: Registering, deploying, and provisioning RHEL images on Google Cloud Managing networking configurations for RHEL Google Cloud instances Configuring platform security and trusted execution technologies Managing Red Hat High Availability (HA) clusters for RHEL instances

## Table of Contents

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION .....	4
<b>CHAPTER 1. INTRODUCING RHEL ON PUBLIC CLOUD PLATFORMS .....</b>	<b>5</b>
1.1. BENEFITS OF USING RHEL IN A PUBLIC CLOUD .....	5
1.2. PUBLIC CLOUD USE CASES FOR RHEL .....	6
1.3. FREQUENT CONCERNS WHEN MIGRATING TO A PUBLIC CLOUD .....	6
1.4. OBTAINING RHEL FOR PUBLIC CLOUD DEPLOYMENTS .....	7
1.5. METHODS FOR CREATING RHEL CLOUD INSTANCES .....	8
<b>CHAPTER 2. PREPARING AND UPLOADING CUSTOM GCE IMAGES TO GOOGLE CLOUD .....</b>	<b>9</b>
2.1. CONFIGURING AND UPLOADING A GCE IMAGE TO GOOGLE CLOUD BY USING THE CLI .....	9
2.2. HOW RHEL IMAGE BUILDER SORTS THE AUTHENTICATION ORDER OF DIFFERENT GOOGLE CLOUD CREDENTIALS .....	10
<b>CHAPTER 3. DEPLOYING A RHEL IMAGE AS A GOOGLE COMPUTE ENGINE INSTANCE ON GOOGLE CLOUD .....</b>	<b>12</b>
3.1. AVAILABLE RHEL IMAGE TYPES FOR PUBLIC CLOUD .....	12
3.2. DEPLOYING A RHEL INSTANCE BY USING A CUSTOM BASE IMAGE .....	13
3.3. UPLOADING AND RUNNING A RHEL INSTANCE ON GOOGLE CLOUD .....	15
3.3.1. Installing the Google Cloud SDK .....	15
3.3.2. Creating a new project on Google Cloud .....	15
3.3.3. Creating and uploading a RHEL image on Google Cloud Storage .....	16
3.3.4. Launching and connecting to a RHEL Google Compute Engine instance .....	17
3.4. ATTACHING RED HAT SUBSCRIPTIONS .....	19
3.5. SETTING UP AUTOMATIC REGISTRATION ON GOOGLE CLOUD GOLD IMAGES .....	19
<b>CHAPTER 4. CONFIGURING A RED HAT HIGH AVAILABILITY CLUSTER ON GOOGLE CLOUD .....</b>	<b>21</b>
4.1. BENEFITS OF USING HIGH-AVAILABILITY CLUSTERS ON PUBLIC CLOUD PLATFORMS .....	21
4.2. REQUIRED SYSTEM PACKAGES .....	22
4.3. CREATING A CUSTOM VIRTUAL PRIVATE CLOUD NETWORK AND SUBNET .....	23
4.4. CREATING AND CONFIGURING A BASE GOOGLE CLOUD INSTANCE .....	23
4.5. CREATING A SNAPSHOT IMAGE .....	26
4.6. CREATING AN HA NODE TEMPLATE INSTANCE AND HA NODES .....	26
4.7. INSTALLING THE HIGH AVAILABILITY PACKAGES AND AGENTS FOR GOOGLE CLOUD .....	27
4.8. CREATING A HIGH AVAILABILITY CLUSTER .....	28
4.9. CONFIGURING FENCING IN A RED HAT HIGH AVAILABILITY CLUSTER .....	29
4.9.1. Displaying available fence agents and their options .....	30
4.9.2. Creating a fence device .....	31
4.9.3. General properties of fencing devices .....	31
4.9.4. Fencing delays .....	38
4.9.5. Testing a fence device .....	40
4.9.6. Configuring fencing levels .....	43
4.9.7. Removing a fence level .....	44
4.9.8. Clearing fence levels .....	44
4.9.9. Verifying nodes and devices in fence levels .....	44
4.9.10. Specifying nodes in fencing topology .....	45
4.9.11. Configuring fencing for redundant power supplies .....	45
4.9.12. Administering fence devices .....	46
4.9.12.1. Displaying configured fence devices .....	46
4.9.12.2. Exporting fence devices as pcs commands .....	46
4.9.12.3. Exporting fence level configuration .....	46
4.9.12.4. Modifying and deleting fence devices .....	47

4.9.12.5. Manually fencing a cluster node	47
4.9.12.6. Disabling a fence device	47
4.9.12.7. Preventing a node from using a fencing device	48
4.9.13. Configuring ACPI for use with integrated fence devices	48
4.9.13.1. Disabling ACPI Soft-Off with the BIOS	49
4.9.13.2. Disabling ACPI Soft-Off in the logind.conf file	50
4.9.13.3. Disabling ACPI completely in the GRUB 2 file	50
4.10. CONFIGURING THE VIRTUAL IP MANAGEMENT RESOURCE AGENT	51
4.10.1. Configuring the primary subnet address range	51
4.10.2. Configuring the secondary subnet address range	52
4.10.3. Additional resources	53
<b>CHAPTER 5. CONFIGURING RHEL ON GOOGLE CLOUD WITH SECURE BOOT .....</b>	<b>54</b>
5.1. TYPES OF RHEL IMAGES ON GOOGLE CLOUD	54
5.2. UNDERSTANDING SECURE BOOT FOR RHEL ON CLOUD	54
5.2.1. Components of Secure Boot	54
5.2.2. Stages of Secure Boot for RHEL on Cloud	55
5.3. CONFIGURING A RHEL INSTANCE FROM A CUSTOM RHEL IMAGE WITH SECURE BOOT	56
5.4. CONFIGURING A RHEL INSTANCE ON THE GOOGLE CLOUD MARKETPLACE WITH SECURE BOOT	58
<b>CHAPTER 6. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH INTEL TDX .....</b>	<b>60</b>
6.1. UNDERSTANDING INTEL TDX SECURE BOOT PROCESS	60
6.2. CONFIGURING A RHEL INSTANCE ON GOOGLE CLOUD WITH INTEL TDX	61
<b>CHAPTER 7. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH AMD SEV SNP .....</b>	<b>63</b>
7.1. PROPERTIES OF SEV-SNP	63
7.2. UNDERSTANDING AMD SEV SNP SECURE BOOT PROCESS	63
7.3. CONFIGURING A RHEL INSTANCE ON GOOGLE CLOUD WITH AMD SEV SNP	64
<b>CHAPTER 8. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH UKI .....</b>	<b>67</b>
8.1. INTRODUCTION TO UNIFIED KERNEL IMAGE	67
8.2. UNDERSTANDING THE UKI SECURE BOOT PROCESS	68



# PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your feedback on our documentation. Let us know how we can improve it.

## Submitting feedback through Jira (account required)

1. Log in to the [Jira](#) website.
2. Click **Create** in the top navigation bar.
3. Enter a descriptive title in the **Summary** field.
4. Enter your suggestion for improvement in the **Description** field. Include links to the relevant parts of the documentation.
5. Click **Create** at the bottom of the dialogue.

# CHAPTER 1. INTRODUCING RHEL ON PUBLIC CLOUD PLATFORMS

Public cloud platforms offer computing resources as a service. Instead of using on-premise hardware, you can run your IT workloads, including Red Hat Enterprise Linux (RHEL) systems, as public cloud instances.

## 1.1. BENEFITS OF USING RHEL IN A PUBLIC CLOUD

Running Red Hat Enterprise Linux (RHEL) on public cloud platforms provides flexible resource allocation, cost efficiency, and software-controlled configurations to optimize your infrastructure without managing physical hardware.

RHEL as a cloud instance located on a public cloud platform has the following benefits over RHEL on-premise physical systems or virtual machines (VMs):

### Flexible and fine-grained allocation of resources

A cloud instance of RHEL runs as a VM on a cloud platform, which means a cluster of remote servers maintained by the cloud service provider. Therefore, on the software level, allocating hardware resources to the instance is easily customizable, such as a specific type of CPU or storage.

In comparison to a local RHEL system, you are also not limited by the capabilities of physical host. Instead, you can select from a variety of features, based on selections offered by the cloud provider.

### Space and cost efficiency

You do not need to own any on-premise servers to host cloud workloads. This avoids the space, power, and maintenance requirements associated with physical hardware.

Instead, on public cloud platforms, you pay the cloud provider directly for using a cloud instance. The cost is typically based on the hardware allocated to the instance and the time to use it. Therefore, you can optimize your costs based on the requirements.

### Software-controlled configurations

You save the entire configuration of a cloud instance as data on the cloud platform and control it with software. Therefore, you can easily create, remove, clone, or migrate the instance. You also operate a cloud instance remotely in a cloud provider console, and it connects to remote storage by default.

In addition, you can back up the current state of a cloud instance as a snapshot at any time.

Afterwards, you can load the snapshot to restore the instance to the saved state.

### Separation from the host and software compatibility

Similarly to a local VM, the RHEL guest operating system on a cloud instance runs on a Kernel-based Virtual Machine (KVM). This kernel is separate from the host operating system and from the *client* system that you use to connect to the instance.

Therefore, you can install any operating system on the cloud instance. This means that on a RHEL public cloud instance, you can run RHEL-specific applications not usable on your local operating system.

In addition, even if the operating system of the instance becomes unstable or compromised, it does not affect your client system.

### Additional resources

- [What is public cloud?](#)

- [What is a hyperscaler?](#)
- [Types of cloud computing](#)

## 1.2. PUBLIC CLOUD USE CASES FOR RHEL

Deploying on a public cloud provides many benefits, but might not be the most efficient solution in every scenario. If you are evaluating whether to migrate your Red Hat Enterprise Linux (RHEL) deployments to the public cloud, consider whether your use case will benefit from the advantages of the public cloud.

### Beneficial use cases

- Deploying public cloud instances is very effective for flexibly increasing and decreasing the active computing power of your deployments. This is also known as *scaling up* and *scaling down*. You can use RHEL on public cloud in the following scenarios:
  - Clusters with high peak workloads and low general performance requirements. Scaling up and down based on your demands can be highly efficient in terms of resource costs.
  - Quickly setting up or expanding your clusters. This avoids high upfront costs of setting up local servers.
- What happens in your local environment does not affect cloud instances. Therefore, you can use them for backup and disaster recovery.

### Potentially problematic use cases

- You are running an existing environment that you cannot adjust. Customizing a cloud instance to fit the specific needs of an existing deployment might not be economically efficient in comparison with your current host platform.
- You are operating with a hard limit on your budget. Maintaining your deployment in a local data center typically provides less flexibility but more control over the maximum resource costs than the public cloud does.

For details on how to obtain RHEL for public cloud deployments, see [Obtaining RHEL for public cloud deployments](#).

### Additional resources

- [Should I migrate my application to the cloud? Here's how to decide.](#)

## 1.3. FREQUENT CONCERNS WHEN MIGRATING TO A PUBLIC CLOUD

Moving your RHEL workloads from a local environment to a public cloud platform might raise concerns about the changes involved. The following are the most commonly asked questions:

### Will my RHEL work differently as a cloud instance than as a local virtual machine?

In most respects, RHEL instances on a public cloud platform work the same as RHEL virtual machines on a local host, such as an on-premise server. Notable exceptions include:

- Instead of private orchestration interfaces, public cloud instances use provider-specific console interfaces for managing your cloud resources.

- Certain features, such as nested virtualization, might not work correctly. If a specific feature is critical for your deployment, check the feature's compatibility in advance with your chosen public cloud provider.

### Will my data stay safe in a public cloud as opposed to a local server?

The data in your RHEL cloud instances is in your ownership, and your public cloud provider does not have any access to it. In addition, major cloud providers support data encryption in transit, which improves the security of data when migrating your virtual machines to the public cloud.

In terms of security of RHEL public cloud instances, the following applies:

- Your public cloud provider is responsible for the security of the cloud hypervisor
- Red Hat provides the security features of the RHEL guest operating systems in your instances
- You manage the specific security settings and practices in your cloud infrastructure

### What effect does my geographic region have on the functionality of RHEL public cloud instances?

You can use RHEL instances on a public cloud platform regardless of your geographical location. Therefore, you can run your instances in the same region as your on-premises server. However, hosting your instances in a physically distant region might cause high latency when operating them. In addition, depending on the public cloud provider, certain regions might offer additional features or be more cost-efficient. Before creating your RHEL instances, review the properties of the hosting regions available for your chosen cloud provider.

## 1.4. OBTAINING RHEL FOR PUBLIC CLOUD DEPLOYMENTS

To deploy Red Hat Enterprise Linux (RHEL) in a public cloud environment, you must select a certified cloud service provider and create a RHEL cloud instance.

- Based on your requirements and the current offerings in the market, select the optimal cloud service provider for your use case. The certified cloud service providers for running RHEL instances are:
  - [Amazon Web Services \(AWS\)](#)
  - [Google Cloud](#)
  - [Microsoft Azure](#)



#### NOTE

This document specifically addresses the process of deploying RHEL on Google Cloud.

- Create a RHEL cloud instance on your cloud platform. For details, see [Methods for creating RHEL cloud instances](#).
- To keep your RHEL deployment up-to-date, use [Red Hat Update Infrastructure \(RHUI\)](#).

#### Additional resources

- [RHUI documentation](#)

- Red Hat Open Hybrid Cloud
- Red Hat Ecosystem Catalog

## 1.5. METHODS FOR CREATING RHEL CLOUD INSTANCES

You can create a Red Hat Enterprise Linux (RHEL) system image and import it to the cloud platform by using either RHEL image builder or purchasing a RHEL image directly from the cloud service provider marketplace. You can then deploy the RHEL image as a cloud instance.

To deploy a RHEL instance on a public cloud platform, you can use either of the following methods:

### Create a RHEL system image and import it to the cloud platform

- To create the system image, you can use the [RHEL image builder](#) or build the image manually.
- This method uses your existing RHEL subscription. This is also referred to as *bring your own subscription* (BYOS).
- You pre-pay a yearly subscription, and you can use your Red Hat customer discount.
- Red Hat provides customer service.
- For creating many images effectively, you can use the **cloud-init** utility.

### Purchase a RHEL instance directly from the cloud provider marketplace

- You post-pay an hourly rate for using the service. This method is also referred to as *pay as you go* (PAYG).
- The cloud service provider provides customer service.

### Additional resources

- [What is a golden image?](#)

# CHAPTER 2. PREPARING AND UPLOADING CUSTOM GCE IMAGES TO GOOGLE CLOUD

With RHEL image builder, you can build a **gce** image, provide credentials for your user or GCP service account, and then upload the **gce** image directly to the GCP environment.

## 2.1. CONFIGURING AND UPLOADING A GCE IMAGE TO GOOGLE CLOUD BY USING THE CLI

Set up a configuration file with credentials to upload your **gce** image to GCP by using the RHEL image builder CLI.



### WARNING

You cannot manually import **gce** image to Google Cloud, because the image will not boot. You must use either **gcloud** or RHEL image builder to upload it.

### Prerequisites

- You have a valid Google account and credentials to upload your image to Google Cloud. The credentials can be from a user account or a service account. The account associated with the credentials must have at least the following IAM roles assigned:
  - **roles/storage.admin** - to create and delete storage objects
  - **roles/compute.storageAdmin** - to import a VM image to Compute Engine.
- You have an existing Google Cloud bucket.

### Procedure

1. Use a text editor to create a **gcp-config.toml** configuration file with the following content:

```
provider = "gcp"
[settings]
bucket = "<gcp_bucket>"
region = "<gcp_storage_region>"
object = "<object_key>"
credentials = "<gcp_credentials>"
```

- **<gcp\_bucket>** points to an existing bucket. It is used to store the intermediate storage object of the image which is being uploaded.
- **<gcp\_storage\_region>** is both a regular Google storage region and a dual or multi-region.
- **<object\_key>** is the name of an intermediate storage object. It must not exist before the upload, and it is deleted when the upload process is done. If the object name does not end with **.tar.gz**, the extension is automatically added to the object name.

- **<gcp\_credentials>** is a **Base64-encoded** scheme of the credentials JSON file downloaded from Google Cloud. The credentials determine which project the Google Cloud uploads the image to.



#### NOTE

Specifying **<gcp\_credentials>** in the **gcp-config.toml** file is optional if you use a different mechanism to authenticate with Google Cloud. For other authentication methods, see [Authenticating with Google Cloud](#).

2. Retrieve the **<gcp\_credentials>** from the JSON file downloaded from Google Cloud.

```
$ sudo base64 -w 0 cee-gcp-nasa-476a1fa485b7.json
```

3. Create a compose with an additional image name and cloud provider profile:

```
$ sudo image-builder build gce --blueprint <blueprint_name> <image_key> gcp-config.toml
```

The image build, upload, and cloud registration processes can take up to ten minutes to complete.

#### Verification

- Verify that the image status is FINISHED:

```
$ sudo image-builder compose status
```

#### Additional resources

- [Identity and Access Management](#)
- [Create storage buckets](#)

## 2.2. HOW RHEL IMAGE BUILDER SORTS THE AUTHENTICATION ORDER OF DIFFERENT GOOGLE CLOUD CREDENTIALS

You can use several different types of credentials with RHEL image builder to authenticate with GCP. If you set RHEL image builder configuration to authenticate with GCP by using multiple sets of credentials, it uses the credentials in an order of preference.

The order of preference is as follows:

1. Credentials specified with the **image-builder** command in the configuration file.
2. **Application Default Credentials** from the **Google Cloud SDK** library, which tries to automatically find a way to authenticate by using the following options:
  - a. If the **GOOGLE\_APPLICATION\_CREDENTIALS** environment variable is set, Application Default Credentials tries to load and use credentials from the file pointed to by the variable.
  - b. Application Default Credentials tries to authenticate by using the service account attached to the resource that is running the code. For example, Google Compute Engine VM.



## NOTE

You must use the Google Cloud credentials to determine which Google Cloud project to upload the image to. Therefore, unless you want to upload all of your images to the same Google Cloud project, you always must specify the credentials in the **gcp-config.toml** configuration file with the **image-builder** command.

# CHAPTER 3. DEPLOYING A RHEL IMAGE AS A GOOGLE COMPUTE ENGINE INSTANCE ON GOOGLE CLOUD

You can deploy a Red Hat Enterprise Linux (RHEL) image as a Google Compute Engine (GCE) virtual machine (VM) by converting it to a GCE-compatible format and then deploying it.

- Use the RHEL image builder. For instructions, see [Preparing and uploading custom GCE images to Google Cloud](#).
- Manually create and configure a [supported format image](#). This is a more complicated process but offers more granular customization options. For details, see the following sections.



## NOTE

For a list of Red Hat product certifications for Google Cloud, see [Red Hat on Google Cloud](#).

To deploy a RHEL image as a GCE instance on Google Cloud, make sure to complete following steps:

- You have created a [Red Hat account](#).
- You have [signed up and set up a Google Cloud account](#).

## 3.1. AVAILABLE RHEL IMAGE TYPES FOR PUBLIC CLOUD

To deploy your RHEL virtual machine VM on a certified cloud service provider (CCSP), you can use several options. The following table lists the available image types, subscriptions, considerations, and sample scenarios for the image types.



## NOTE

To deploy customized ISO images, you can use RHEL image builder. With RHEL image builder, you can create, upload, and deploy these custom images specific to your chosen CCSP. For details, see [Composing a Customized RHEL System Image](#).

**Table 3.1. Image options**

Image types	Subscriptions	Considerations	Sample scenario
Deploy a Red Hat gold image	Use your existing Red Hat subscriptions	The subscriptions include the Red Hat product cost and support for Cloud Access images, while you pay the CCSP for all other instance costs	Select a Red Hat gold image on the CCSP. For details on gold images and how to access them on the CCSP, see the <a href="#">Red Hat Cloud Access Reference Guide</a>

Image types	Subscriptions	Considerations	Sample scenario
Deploy a custom image that you move to the CCSP	Use your existing Red Hat subscriptions	The subscriptions includes the Red Hat product cost and support for custom RHEL image, while you pay the CCSP for all other instance costs	Upload your custom image and attach your subscriptions
Deploy an existing RHEL based custom machine image	The custom machine images include a RHEL image	You pay the CCSP on an hourly basis based on a <i>pay-as-you-go</i> model. For this model, on-demand images are available on the CCSP marketplace. The CCSP provides support for these images, while Red Hat handles updates. The CCSP provides updates through the Red Hat Update Infrastructure (RHUI)	Select a RHEL image when you launch an instance on the CCSP cloud management console, or choose an image from the CCSP marketplace.



## IMPORTANT

You cannot convert an on-demand instance to a custom RHEL instance. For migrating from an on-demand image to a custom RHEL bring your own subscription (BYOS) image, do the following:

- Create a new custom RHEL instance, then migrate data from your on-demand instance.
- When you complete data migration, terminate the on-demand instance to avoid additional billing.

## Additional resources

- [Red Hat Ecosystem Catalog](#)
- [Red Hat Cloud Access Reference Guide](#)
- [Compute Engine images](#)
- [Create and start a Compute Engine instance](#)

## 3.2. DEPLOYING A RHEL INSTANCE BY USING A CUSTOM BASE IMAGE

To manually configure a virtual machine (VM), first create a base (starter) image. Then, you can modify configuration settings and add the packages the VM requires to operate on the cloud. You can also make additional configuration changes for your specific application after you upload the image.

Creating a VM from a base image has the following advantages:

- Fully customizable
- High flexibility for any use case
- Lightweight - includes only the operating system and the required runtime libraries

To create a custom base image of RHEL from an ISO image, you can use the command line interface (CLI) or the web console for creating and configuring VM.



### NOTE

Verify the following VM configurations.

- SSH - Enable SSH to give remote access to your VM.
- DHCP - Configure the primary virtual adapter to use DHCP.

### Prerequisites

- You have [enabled virtualization](#) on the host machine.
- For web console, ensure the following options:
  - You have not checked the **Immediately Start VM** option.
  - You have already changed the **Memory** size to your preferred settings.
- You have changed the **Model** option under **Virtual Network Interface Settings** to **virtio** and **vCPUs** to the capacity settings for the VM.

### Procedure

1. Configure the Red Hat Enterprise Linux (RHEL) VM:
  - a. To install from the command line (CLI), ensure that you set the default memory, network interfaces, and CPUs according to your requirement for the VM. For details, see [Creating virtual machines by using the command line](#)
  - b. To install from the web console, see [Creating virtual machines by using the web console](#)
2. When the installation starts:
  - a. Create a **root** password.
  - b. Create an administrative user account.
3. After the installation completes, reboot the VM and log in to the **root** account.
4. After logging in as **root**, you can configure the image.
5. Register the VM and enable the RHEL repository:

```
# subscription-manager register
```

## Verification

- Verify if the system has the **cloud-init** package and enable it:

```
# dnf install cloud-init  
# systemctl enable --now cloud-init.service
```

- Power off the VM.

## 3.3. UPLOADING AND RUNNING A RHEL INSTANCE ON GOOGLE CLOUD

To run your RHEL instance on Google Cloud, you need to configure and upload a RHEL image to Google Cloud.

### 3.3.1. Installing the Google Cloud SDK

You can use Google Cloud SDK to manage Google Cloud resources and services from your command line by using Google Cloud CLI.

#### Prerequisites

- You have downloaded, extracted, and initialized [the Google Cloud SDK](#).

#### Procedure

- Use the **gcloud** CLI utility to manage project and instances.
- Create a project:

```
# gcloud projects create <example_gcp_project_id> --name <example_gcp_project>
```

The example creates a project with the project ID **<example\_gcp\_project\_id>** and the project name **<example\_gcp\_project>**.

- Display project information:

```
# gcloud compute <example_project_info> describe --project <example_project_name>
```

#### Additional resources

- [gcloud project create](#)
- [gcloud command-line tool overview](#)

### 3.3.2. Creating a new project on Google Cloud

To upload your RHEL image to Google Cloud, You need to create new project on Google Cloud. A project manages your assigned Google Cloud resources.

## Prerequisites

- You have an account on [Google Cloud](#).

## Procedure

1. Launch the [Google Cloud Console](#).
2. Click the drop-down menu to the right of **Google Cloud**.
3. From the pop-up menu, click **NEW PROJECT**.
4. From the **New Project** window, enter a name for your new project.
5. Check **Organization**. Click the drop-down menu to change the organization, if necessary.
6. Confirm the **Location** of your parent organization or folder. Click **Browse** to search for and change this value, if necessary.
7. Click **CREATE** to create your new Google Cloud project.

## Additional resources

- [Creating and Managing Resources in Google Cloud](#)

### 3.3.3. Creating and uploading a RHEL image on Google Cloud Storage

To import and store objects such as virtual machine (VM) images, you can create a Google Cloud storage bucket.

## Procedure

1. Log in to the Google cloud console:

```
# gcloud auth login
```

2. Create a storage bucket:

```
# gsutil mb gs://<example_bucket_name>
```



#### NOTE

Alternatively, you can use the Google Cloud Console to create a bucket. For details, see [Create a bucket](#).

3. Specify the image name that you want to create, the existing bucket name, and the name of the image:

```
# *gcloud compute images create my-image-name --source-uri  
gs://__<example_bucket_name>__/disk.raw.tar.gz*
```

**NOTE**

Alternatively, you can use the Google Cloud Console to create an image. See [Creating, deleting, and deprecating custom images](#) for more information.

4. Optional: Find the image in the Google Cloud Console.
  - a. Click the **Navigation** menu to the left of the **Google Cloud Console** banner.
  - b. Select **Compute Engine** and then **Images**.
5. Run the **qemu-img** command to convert your **qcow2** image to the **raw** format:

```
# qemu-img convert -f qcow2 -O raw rhel-10.0-sample.qcow2 disk.raw
```

6. Compress the image:

```
# tar --format=oldgnu -Sczf disk.raw.tar.gz disk.raw
```

7. Upload the image to the existing bucket:

```
# gsutil cp disk.raw.tar.gz gs://<example_bucket_name>
```

**NOTE**

Upload could take a few minutes.

8. From the **Google Cloud** home screen, click the collapsed menu icon and select **Storage** and then select **Browser**.
9. Click the name of your bucket where **disk.raw.tar.gz** is now listed.

**NOTE**

You can also upload your image by using the **Google Cloud Console**. To do so, click the bucket name and then click **Upload files**.

## Additional resources

- [gcloud compute images create](#)
- [Manually importing virtual disks](#)
- [Choosing an import method](#)
- [Installing the gsutil command-line tool](#)

### 3.3.4. Launching and connecting to a RHEL Google Compute Engine instance

You can configure a Google Compute Engine (GCE) virtual machine (VM) instance from an image by using the Google Cloud Console.

#### Procedure

1. From the [Google Cloud Console Dashboard page](#), click the **Navigation** menu to the left of the **Google Cloud Console banner** and select **Compute Engine** and then select **Images**.
2. Select your image.
3. Click **Create Instance**.
4. On the **Create an instance** page, enter a **Name** for your instance.
5. Choose a **Region** and **Zone**.
6. Choose a **Machine configuration** that meets or exceeds the requirements of your workload.
7. Ensure that **Boot disk** specifies the name of your image.
8. Optional: Under **Firewall**, select **Allow HTTP traffic** or **Allow HTTPS traffic**.
9. Click **Create**.
10. Find your image under **VM instances**.
11. Click the **Navigation** menu to the left of the **Google Cloud Console banner** and select **Compute Engine** and then select **VM instances**.



#### NOTE

Alternatively, you can use the **gcloud compute instances create** command to create a GCE VM instance from an image.

```
# gcloud compute instances create myinstance3 --zone=us-central1-a --image test-iso2-image
```

The example creates a VM instance named **myinstance3** in zone **us-central1-a** based upon the existing image **test-iso2-image**. For details, see [gcloud compute instances create](#).

12. Use the **ssh-keygen** utility to generate an SSH key pair to use with GCE by using the public IP address:

```
# ssh-keygen -t rsa -f ~/.ssh/google_compute_engine.pub
```
13. From the [Google Cloud Console Dashboard page](#), click the **Navigation** menu to the left of the **Google Cloud Console banner** and select **Compute Engine** and then select **Metadata**.
14. Click **SSH Keys** and then click **Edit**.
15. Enter the output generated from the **~/.ssh/google\_compute\_engine.pub** file and click **Save**.
16. Connect to the instance:

```
# ssh -i ~/.ssh/google_compute_engine <username>@<instance_external_ip>
```



## NOTE

Or, you can run the **gcloud compute config-ssh** command to populate the config file with aliases for instances. The aliases allow simple SSH connections by instance name. For details, see [gcloud compute config-ssh](#).

### Additional resources

- [Connecting to instances](#)
- [Creating and starting a VM instance](#)

## 3.4. ATTACHING RED HAT SUBSCRIPTIONS

To register and attach your Red Hat subscription to a RHEL instance, you can use the **subscription-manager** command.

### Prerequisites

- You have an active [Red Hat account](#).

### Procedure

1. Register your system:

```
# subscription-manager register
```

2. Attach your subscriptions:

- You can use an activation key to attach subscriptions. See [Creating Red Hat Customer Portal Activation Keys](#) for more information.
- Otherwise, you can manually attach a subscription by using the ID of the subscription pool (Pool ID). See [Attaching a host-based subscription to hypervisors](#).

3. Optional: To collect various system metrics about the instance in the [Red Hat Hybrid Cloud Console](#), you can register the instance with [Red Hat Lightspeed](#).

```
# insights-client register --display-name <display_name_value>
```

For information about further configuration of Red Hat Lightspeed, see [Client Configuration Guide for Red Hat Lightspeed](#).

### Additional resources

- [Creating Red Hat Customer Portal Activation Keys](#)
- [Client Configuration Guide for Red Hat Lightspeed](#)
- [Red Hat Cloud Access Reference Guide](#)

## 3.5. SETTING UP AUTOMATIC REGISTRATION ON GOOGLE CLOUD GOLD IMAGES

You can deploy Red Hat Enterprise Linux (RHEL) virtual machines (VMs) more efficiently on Google Cloud by using gold images of RHEL. This ensures that the VMs are automatically registered to the Red Hat Subscription Manager (RHSM).

## Prerequisites

- RHEL gold images are available to you in Google Cloud. For instructions, see [Using gold images on Google](#).



### NOTE

A Google Cloud account can only be attached to a single Red Hat account at a time. Therefore, ensure no other users require access to the Google Cloud account before attaching it to your Red Hat account.

## Procedure

1. Use the gold image to create a RHEL VM in your Google Cloud instance. For instructions, see [Uploading and running a RHEL instance on Google Cloud](#).  
If your RHSM settings are correct, the VM will be automatically subscribed to RHSM.

## Verification

- In a RHEL VM created by using the above instructions, verify the system is registered to RHSM. On a successfully registered system, the `subscription-manager identity` command displays the UUID of the system. For example:

```
# subscription-manager identity
system identity: fdc46662-c536-43fb-a18a-bbcb283102b7
name: 192.168.122.222
org name: 6340056
org ID: 6340056
```

## Additional resources

- [Red Hat gold images in Azure](#)
- [Overview of RHEL images in Azure](#)
- [Configuring cloud integration for Red Hat services](#)

# CHAPTER 4. CONFIGURING A RED HAT HIGH AVAILABILITY CLUSTER ON GOOGLE CLOUD

To group RHEL nodes on Google Cloud and automatically redistribute workloads if a node fails, you can configure a high availability (HA) cluster on Google Cloud. The process for setting up HA clusters on Google Cloud is comparable to configuring them in traditional, non-cloud environments.

To configure a Red Hat HA cluster on Google Cloud that uses Google Compute Engine (GCE) instances as cluster nodes, see the following sections. You have several options for obtaining RHEL images for the cluster. For details, see [Available RHEL image types for public cloud](#).

Before you begin, make sure to complete following steps:

- You have created a Red Hat account. For details, see [Red Hat account](#).
- You have signed up and set up a Google Cloud account. For details, see [Google Cloud account](#).
- Red Hat Enterprise Linux 10 Server: **rhel-10-server-rpms/8Server/x86\_64**
- Red Hat Enterprise Linux 10 Server (High Availability): **rhel-10-server-ha-rpms/8Server/x86\_64**
- Your project should have a service account that belongs to a VM instance and not an individual user. For details, see [Service accounts](#).
- You can use the default service account instead of creating a separate service account. For details, see [Using the Compute Engine Default Service Account](#).

If you or your project administrator create a custom service account, configure the service account for the following roles:

- Cloud Trace Agent
- Compute Admin
- Compute Network Admin
- Cloud Datastore User
- Logging Admin
- Monitoring Editor
- Monitoring Metric Writer
- Service Account Administrator
- Storage Admin

## 4.1. BENEFITS OF USING HIGH-AVAILABILITY CLUSTERS ON PUBLIC CLOUD PLATFORMS

A high-availability (HA) cluster links a set of computers (called *nodes*) to run a specific workload. HA clusters offer redundancy to handle hardware or software failures. When a node in the HA cluster fails, the Pacemaker cluster resource manager quickly distributes the workload to other nodes, ensuring that services on the cluster continue without noticeable downtime.

You can also run HA clusters on public cloud platforms. In this case, you would use virtual machine (VM) instances in the cloud as the individual cluster nodes. Using HA clusters on a public cloud platform has the following benefits:

- *Improved availability*: In case of a VM failure, the workload is quickly redistributed to other nodes, so running services are not disrupted.
- *Scalability*: You can start additional nodes when demand is high and stop them when demand is low.
- *Cost-effectiveness*: With the pay-as-you-go pricing, you pay only for nodes that are running.
- *Simplified management*: Some public cloud platforms offer management interfaces to make configuring HA clusters easier.

To enable HA on your RHEL systems, Red Hat offers a HA Add-On. You can configure a RHEL cluster with Red Hat HA Add-On to manage HA clusters with groups of RHEL servers. Red Hat HA Add-On gives access to integrated and streamlined tools. With cluster resource manager, fencing agents, and resource agents, you can set up and configure the cluster for automation. The Red Hat HA Add-On offers the following components for automation:

- **Pacemaker**, a cluster resource manager that offers both a command line utility (**pcs**) and a GUI (**pcsd**) to support many nodes
- **Corosync** and **Kronosnet** to create and manage HA clusters
- Resource agents to configure and manage custom applications
- Fencing agents to use cluster on platforms such as bare-metal servers and virtual machines

The Red Hat HA Add-On handles critical tasks such as node failures, load balancing, and node health checks for fault tolerance and system reliability.

#### Additional resources

- [High Availability Add-On overview](#)

## 4.2. REQUIRED SYSTEM PACKAGES

To create and configure a base image of RHEL, your host system must have the following packages installed.

**Table 4.1. System packages**

Package	Repository	Description
libvirt	<b>rhel-10-for-x86_64-appstream-rpms</b>	Open source API, daemon, and management tool for managing platform virtualization
virt-install	<b>rhel-10-for-x86_64-appstream-rpms</b>	A command-line utility for building VMs

Package	Repository	Description
libguestfs	<b>rhel-10-for-x86_64-appstream-rpms</b>	A library for accessing and modifying VM file systems
guestfs-tools	<b>rhel-10-for-x86_64-appstream-rpms</b>	System administration tools for VMs; includes the <b>virt-customize</b> utility

By confirming installation of the mentioned system packages, you can now follow the deployment steps in [Deploying a RHEL instance by using a custom base image](#)

## 4.3. CREATING A CUSTOM VIRTUAL PRIVATE CLOUD NETWORK AND SUBNET

You must configure a custom virtual private cloud (VPC) network and subnet to enable High Availability (HA) for your cluster.

### Prerequisites

- You have installed [the Google Cloud SDK](#).
- You have created [a new Google Cloud project](#).

### Procedure

1. Launch the Google Cloud Console.
2. Select **VPC networks** under **Networking** in the left navigation pane.
3. Click **Create VPC Network**
4. Enter a name for the VPC network.
5. Under the **New subnet**, create a **Custom subnet** in the region where you want to create the cluster.
6. Click **Create**.

## 4.4. CREATING AND CONFIGURING A BASE GOOGLE CLOUD INSTANCE

Create and configure a Google Cloud instance that complies with the platform's operating and security standards. Use this foundational instance as a starting point to deploy applications, apply additional configurations, or set up clusters as needed.

### Procedure

1. Create an image from the compressed file in the bucket.

```
$ gcloud compute images create BaseImageName --source-uri
gs://BucketName/BaseImageName.tar.gz
```

Example:

```
$ gcloud compute images create rhel-76-server --source-uri gs://user-rhelha/rhel-server-
76.tar.gz
Created [https://www.googleapis.com/compute/v1/projects/MyProject/global/images/rhel-
server-76].
NAME      PROJECT      FAMILY DEPRECATED STATUS
rhel-76-server  rhel-ha-testing-on-gcp          READY
```

2. Create a template instance from the image. The minimum size required for a base RHEL instance is n1-standard-2. See [gcloud compute instances create](#) for additional configuration options.

```
$ gcloud compute instances create BaseInstanceName --can-ip-forward --machine-type n1-
standard-2 --image BaseImageName --service-account ServiceAccountEmail
```

Example:

```
$ gcloud compute instances create rhel-76-server-base-instance --can-ip-forward --machine-
type n1-standard-2 --image rhel-76-server --service-account account@project-name-on-
gcp.iam.gserviceaccount.com
Created [https://www.googleapis.com/compute/v1/projects/rhel-ha-testing-on-gcp/zones/us-
east1-b/instances/rhel-76-server-base-instance].
NAME ZONE MACHINE_TYPE PREEMPTIBLE INTERNAL_IP EXTERNAL_IP
STATUS
rhel-76-server-base-instance us-east1-bn1-standard-2      10.10.10.3  192.227.54.211
RUNNING
```

3. Connect to the instance with an SSH terminal session.

```
$ ssh root@PublicIPaddress
```

4. Update the RHEL software.

- Register with Red Hat Subscription Manager (RHSM).
- Enable a Subscription Pool ID.
- Disable all repositories.

```
# subscription-manager repos --disable=
```

- Enable the following repository.

```
# subscription-manager repos --enable=rhel-10-server-rpms
```

- Run the **dnf update** command.

```
# dnf update -y
```

5. Install the Google Cloud Linux Guest Environment on the running instance (in-place installation).  
See [Install the guest environment in-place](#) for instructions.
6. Select the **CentOS/RHEL** option.
7. Copy the command script and paste it to the terminal to run the script immediately.
8. Make the following configuration changes to the instance. These changes are based on Google Cloud recommendations for custom images. See [gcloudcompute images list](#) for more information.
  - a. Edit the **/etc/chrony.conf** file and remove all NTP servers.
  - b. Add the following NTP server.

```
metadata.google.internal iburst Google NTP server
```

- c. Remove any persistent network device rules.

```
# rm -f /etc/udev/rules.d/70-persistent-net.rules  
# rm -f /etc/udev/rules.d/75-persistent-net-generator.rules
```

- d. Set the network service to start automatically.

```
# chkconfig network on
```

- e. Set the **sshd** service to start automatically.

```
# systemctl enable sshd  
# systemctl is-enabled sshd
```

- f. Set the time zone to UTC.

```
# ln -sf /usr/share/zoneinfo/UTC /etc/localtime
```

- g. Optional: Edit the **/etc/ssh/ssh\_config** file and add the following lines to the end of the file. This keeps your SSH session active during longer periods of inactivity.

```
# Server times out connections after several minutes of inactivity.  
# Keep alive ssh connections by sending a packet every 7 minutes.  
ServerAliveInterval 420
```

- h. Edit the **/etc/ssh/sshd\_config** file and make the following changes, if necessary. The **ClientAliveInterval 420** setting is optional; this keeps your SSH session active during longer periods of inactivity.

```
PermitRootLogin no  
PasswordAuthentication no  
AllowTcpForwarding yes  
X11Forwarding no  
PermitTunnel no
```

```
# Compute times out connections after 10 minutes of inactivity.  
# Keep ssh connections alive by sending a packet every 7 minutes.  
ClientAliveInterval 420
```

9. Disable password access.

```
ssh_pwauth from 1 to 0.  
ssh_pwauth: 0
```



### IMPORTANT

Previously, you enabled password access to allow SSH session access to configure the instance. You must disable password access. All SSH session access must be without password.

10. Unregister the instance from the subscription manager.

```
# subscription-manager unregister
```

11. Clean the shell history. Keep the instance running for the next procedure.

```
# export HISTSIZE=0
```

## 4.5. CREATING A SNAPSHOT IMAGE

To preserve the configuration and disk data of a Google Cloud high availability (HA) instance, you can create a snapshot of the instance. A snapshot provides a point-in-time backup of your disk, allowing you to restore your system or create new instances with the same configuration and disk data.

### Procedure

1. On the running instance, synchronize data to disk.

```
# sync
```

2. On your host system, create the snapshot.

```
$ gcloud compute disks snapshot InstanceName --snapshot-names SnapshotName
```

3. On your host system, create the configured image from the snapshot.

```
$ gcloud compute images create ConfiguredImageFromSnapshot --source-snapshot  
SnapshotName
```

### Additional resources

- [Creating Persistent Disk Snapshots](#)

## 4.6. CREATING AN HA NODE TEMPLATE INSTANCE AND HA NODES

You can create a node template from a configured snapshot image, then use this template to create all HA nodes.

### Procedure

1. Create an instance template:

```
$ gcloud compute instance-templates create InstanceTemplateName --can-ip-forward --machine-type n1-standard-2 --image ConfiguredImageFromSnapshot --service-account ServiceAccountEmailAddress
```

Example:

```
$ gcloud compute instance-templates create rhel-10-instance-template --can-ip-forward --machine-type n1-standard-2 --image rhel-10-gcp-image --service-account account@project-name-on-gcp.iam.gserviceaccount.com
Created [https://www.googleapis.com/compute/v1/projects/project-name-on-gcp/global/instanceTemplates/rhel-10-instance-template].
NAME MACHINE_TYPE PREEMPTIBLE CREATION_TIMESTAMP
rhel-101-instance-template n1-standard-2 2018-07-25T11:09:30.506-07:00
```

2. Create multiple nodes in one zone:

```
# gcloud compute instances create NodeName01 NodeName02 --source-instance-template InstanceTemplateName --zone RegionZone --network=NetworkName --subnet=SubnetName
```

Example:

```
$ gcloud compute instances create rhel10-node-01 rhel10-node-02 rhel10-node-03 --source-instance-template rhel-10-instance-template --zone us-west1-b --network=projectVPC --subnet=range0
Created [https://www.googleapis.com/compute/v1/projects/project-name-on-gcp/zones/us-west1-b/instances/rhel81-node-01].
Created [https://www.googleapis.com/compute/v1/projects/project-name-on-gcp/zones/us-west1-b/instances/rhel81-node-02].
Created [https://www.googleapis.com/compute/v1/projects/project-name-on-gcp/zones/us-west1-b/instances/rhel81-node-03].
NAME ZONE MACHINE_TYPE PREEMPTIBLE INTERNAL_IP EXTERNAL_IP
STATUS
rhel81-node-01 us-west1-b n1-standard-2 10.10.10.4 192.230.25.81 RUNNING
rhel81-node-02 us-west1-b n1-standard-2 10.10.10.5 192.230.81.253 RUNNING
rhel81-node-03 us-east1-b n1-standard-2 10.10.10.6 192.230.102.15 RUNNING
```

## 4.7. INSTALLING THE HIGH AVAILABILITY PACKAGES AND AGENTS FOR GOOGLE CLOUD

To configure a Red Hat High Availability (HA) cluster on Google Cloud, you must install the High Availability packages and agents on each node.

### Procedure

1. In the Google Cloud Console, select **Compute Engine** and then select **VM instances**.

2. Select the instance, click the arrow next to **SSH**, and select the **View gcloud command option**.
3. Paste this command at the terminal to access the instance without password.
4. Enable sudo account access and register with Red Hat Subscription Manager.
5. Enable a Subscription Pool ID.
6. Disable all repositories.

```
# subscription-manager repos --disable= *
```

7. Enable the following repositories.

```
# subscription-manager repos --enable=rhel-10-server-rpms  
# subscription-manager repos --enable=rhel-10-for-x86_64-highavailability-rpms
```

8. Install **pcs pacemaker**, the fence agents, and the resource agents.

```
# dnf install -y pcs pacemaker fence-agents-gce resource-agents-gcp
```

9. Update all packages.

```
# dnf update -y
```

## 4.8. CREATING A HIGH AVAILABILITY CLUSTER

You can create a Red Hat High Availability Add-On cluster. This example uses nodes **z1.example.com** and **z2.example.com**.



### NOTE

To display the parameters of a **pcs** command and a description of those parameters, use the **-h** option of the **pcs** command.

#### Prerequisites

- You have created a [Red Hat account](#)
- You have [signed up and set up an Azure account](#).

#### Procedure

1. Authenticate the **pcs** user **hacluster** for each node in the cluster on the node from which you will be running **pcs**.

The following command authenticates user **hacluster** on **z1.example.com** for both of the nodes in a two-node cluster that will consist of **z1.example.com** and **z2.example.com**.

```
[root@z1 ~]# pcs host auth z1.example.com z2.example.com  
Username: hacluster  
Password:
```

**z1.example.com:** Authorized  
**z2.example.com:** Authorized

2. Execute the following command from **z1.example.com** to create the two-node cluster **my\_cluster** that consists of nodes **z1.example.com** and **z2.example.com**. This will propagate the cluster configuration files to both nodes in the cluster. This command includes the **--start** option, which will start the cluster services on both nodes in the cluster.

```
[root@z1 ~]# pcs cluster setup my_cluster --start z1.example.com z2.example.com
```

3. Enable the cluster services to run on each node in the cluster when the node is booted.



#### NOTE

For your particular environment, you can skip this step by keeping the cluster services disabled. If enabled and a node goes down, any issues with your cluster or your resources are resolved before the node rejoins the cluster. If you keep the cluster services disabled, you need to manually start the services when you reboot a node by executing the **pcs cluster start** command on that node.

```
[root@z1 ~]# pcs cluster enable --all
```

4. Display the status of the cluster you created with the **pcs cluster status** command. Because there could be a slight delay before the cluster is up and running when you start the cluster services with the **--start** option of the **pcs cluster setup** command, you should ensure that the cluster is up and running before performing any subsequent actions on the cluster and its configuration.

```
[root@z1 ~]# pcs cluster status
Cluster Status:
Stack: corosync
Current DC: z2.example.com (version 2.0.0-10.el8-b67d8d0de9) - partition with quorum
Last updated: Thu Oct 11 16:11:18 2018
Last change: Thu Oct 11 16:11:00 2018 by hacluster via crmd on z2.example.com
2 Nodes configured
0 Resources configured

...
```

## 4.9. CONFIGURING FENCING IN A RED HAT HIGH AVAILABILITY CLUSTER

When a node becomes unresponsive, the cluster must isolate it to prevent data corruption. Since the node cannot be contacted directly, you must configure fencing. An external fence device cuts off the node's access to shared resources or performs a hard reboot.

Without a fence device configured you do not have a way to know that the resources previously used by the disconnected cluster node have been released, and this could prevent the services from running on any of the other cluster nodes. Conversely, the system may assume erroneously that the cluster node has released its resources and this can lead to data corruption and data loss. Without a fence device configured data integrity cannot be guaranteed and the cluster configuration will be unsupported.

When the fencing is in progress no other cluster operation is allowed to run. Normal operation of the cluster cannot resume until fencing has completed or the cluster node rejoins the cluster after the cluster node has been rebooted. For more information about fencing and its importance in a Red Hat High Availability cluster, see the Red Hat Knowledgebase solution [Fencing in a Red Hat High Availability Cluster](#).

#### 4.9.1. Displaying available fence agents and their options

You can view available fencing agents and the available options for specific fencing agents.



##### NOTE

Your system's hardware determines the type of fencing device to use for your cluster. For information about supported platforms and architectures and the different fencing devices, see the Red Hat Knowledgebase article [Cluster Platforms and Architectures](#) section of the article [Support Policies for RHEL High Availability Clusters](#) .

Run the following command to list all available fencing agents. When you specify a filter, this command displays only the fencing agents that match the filter.

```
# pcs stonith list [filter]
```

Run the following command to display the options for the specified fencing agent.

```
# pcs stonith describe [stonith_agent]
```

For example, the following command displays the options for the fence agent for APC over telnet/SSH.

```
# pcs stonith describe fence_apc
Stonith options for: fence_apc
ipaddr (required): IP Address or Hostname
login (required): Login Name
passwd: Login password or passphrase
passwd_script: Script to retrieve password
cmd_prompt: Force command prompt
secure: SSH connection
port (required): Physical plug number or name of virtual machine
identity_file: Identity file for ssh
switch: Physical switch number on device
inet4_only: Forces agent to use IPv4 addresses only
inet6_only: Forces agent to use IPv6 addresses only
ipport: TCP port to use for connection with device
action (required): Fencing Action
verbose: Verbose mode
debug: Write debug information to given file
version: Display version information and exit
help: Display help and exit
separator: Separator for CSV created by operation list
power_timeout: Test X seconds for status change after ON/OFF
shell_timeout: Wait X seconds for cmd prompt after issuing command
login_timeout: Wait X seconds for cmd prompt after login
power_wait: Wait X seconds after issuing ON/OFF
delay: Wait X seconds before fencing is started
retry_on: Count of attempts to retry power on
```

**WARNING**

For fence agents that provide a **method** option, with the exception of the **fence\_sbd** agent a value of **cycle** is unsupported and should not be specified, as it may cause data corruption. Even for **fence\_sbd**, however, you should not specify a method and instead use the default value.

### 4.9.2. Creating a fence device

Create a fence device using the **pcs stonith create** command. To view all available creation options, use the **pcs stonith -h** command.

#### Procedure

- Create a fence device:

```
# pcs stonith create stonith_id stonith_device_type [stonith_device_options] [op operation_action operation_options]
```

- The following command creates a single fencing device for a single node:

```
# pcs stonith create MyStonith fence_virt pcmk_host_list=f1 op monitor interval=30s
```

Some fence devices can fence only a single node, while other devices can fence multiple nodes. The parameters you specify when you create a fencing device depend on what your fencing device supports and requires.

- Some fence devices can automatically determine what nodes they can fence.
- You can use the **pcmk\_host\_list** parameter when creating a fencing device to specify all of the machines that are controlled by that fencing device.
- Some fence devices require a mapping of host names to the specifications that the fence device understands. You can map host names with the **pcmk\_host\_map** parameter when creating a fencing device.

#### Additional resources

- [General properties of fencing devices](#).
- [Testing a fence device](#).

### 4.9.3. General properties of fencing devices

Configure fencing behavior using device-specific options and cluster-wide properties. Device options define agent settings, such as IP addresses, and metadata like delays. Cluster properties manage global logic, including timeouts and the **stonith-enabled** parameter.

Any cluster node can fence any other cluster node with any fence device, regardless of whether the fence resource is started or stopped. Whether the resource is started controls only the recurring monitor for the device, not whether it can be used, with the following exceptions:

- You can disable a fencing device by running the **pcs stonith disable *stonith\_id*** command. This will prevent any node from using that device.
- To prevent a specific node from using a fencing device, you can configure location constraints for the fencing resource with the **pcs constraint location ... avoids** command.
- Configuring **stonith-enabled=false** will disable fencing altogether. Note, however, that Red Hat does not support clusters when fencing is disabled, as it is not suitable for a production environment.

The following table describes the general properties you can set for fencing devices.

**Table 4.2. General Properties of Fencing Devices**

Field	Type	Default	Description
<b>pcmk_host_map</b>	string		A mapping of host names to port numbers for devices that do not support host names. For example: <b>node1:1;node2:2,3</b> tells the cluster to use port 1 for node1 and ports 2 and 3 for node2. The <b>pcmk_host_map</b> property supports special characters inside <b>pcmk_host_map</b> values using a backslash in front of the value. For example, you can specify <b>pcmk_host_map="node3:plug\1"</b> to include a space in the host alias.
<b>pcmk_host_list</b>	string		A list of machines controlled by this device (Optional unless <b>pcmk_host_check=static-list</b> ).
<b>pcmk_host_check</b>	string	<ul style="list-style-type: none"> <li>* <b>static-list</b> if either <b>pcmk_host_list</b> or <b>pcmk_host_map</b> is set</li> <li>* Otherwise, <b>dynamic-list</b> if the fence device supports the <b>list</b> action</li> <li>* Otherwise, <b>status</b> if the fence device supports the <b>status</b> action</li> <li>*Otherwise, <b>none</b>.</li> </ul>	How to determine which machines are controlled by the device. Allowed values: <b>dynamic-list</b> (query the device), <b>static-list</b> (check the <b>pcmk_host_list</b> attribute), <b>none</b> (assume every device can fence every machine)

The following table summarizes additional properties you can set for fencing devices. Note that these properties are for advanced use only.

**Table 4.3. Advanced Properties of Fencing Devices**

Field	Type	Default	Description
<b>pcmk_host_argument</b>	string	port	An alternate parameter to supply instead of port. Some devices do not support the standard port parameter or may provide additional ones. Use this to specify an alternate, device-specific parameter that should indicate the machine to be fenced. A value of <b>none</b> can be used to tell the cluster not to supply any additional parameters.
<b>pcmk_reboot_action</b>	string	reboot	An alternate command to run instead of <b>reboot</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the reboot action.
<b>pcmk_reboot_timeout</b>	time	60s	Specify an alternate timeout to use for reboot actions instead of <b>stonith-timeout</b> . Some devices need much more/less time to complete than normal. Use this to specify an alternate, device-specific, timeout for reboot actions.
<b>pcmk_reboot_retries</b>	integer	2	The maximum number of times to retry the <b>reboot</b> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries reboot actions before giving up.
<b>pcmk_off_action</b>	string	off	An alternate command to run instead of <b>off</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the off action.

Field	Type	Default	Description
<b>pcmk_off_timeout</b>	time	60s	Specify an alternate timeout to use for off actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for off actions.
<b>pcmk_off_retries</b>	integer	2	The maximum number of times to retry the off command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries off actions before giving up.
<b>pcmk_list_action</b>	string	list	An alternate command to run instead of <b>list</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the list action.
<b>pcmk_list_timeout</b>	time	60s	Specify an alternate timeout to use for list actions. Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for list actions.
<b>pcmk_list_retries</b>	integer	2	The maximum number of times to retry the <b>list</b> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries list actions before giving up.
<b>pcmk_monitor_action</b>	string	monitor	An alternate command to run instead of <b>monitor</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the monitor action.

Field	Type	Default	Description
<b>pcmk_monitor_timeout</b>	time	60s	Specify an alternate timeout to use for monitor actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for monitor actions.
<b>pcmk_monitor_retries</b>	integer	2	The maximum number of times to retry the <b>monitor</b> command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries monitor actions before giving up.
<b>pcmk_status_action</b>	string	status	An alternate command to run instead of <b>status</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the status action.
<b>pcmk_status_timeout</b>	time	60s	Specify an alternate timeout to use for status actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for status actions.
<b>pcmk_status_retries</b>	integer	2	The maximum number of times to retry the status command within the timeout period. Some devices do not support multiple connections. Operations may fail if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries status actions before giving up.
<b>pcmk_delay_base</b>	string	0s	Enables a base delay for fencing actions and specifies a base delay value. You can specify different values for different nodes with the <b>pcmk_delay_base</b> parameter. For general information about fencing delay parameters and their interactions, see <a href="#">Fencing delays</a> .

Field	Type	Default	Description
<b>pcmk_delay_max</b>	time	0s	Enables a random delay for fencing actions and specifies the maximum delay, which is the maximum value of the combined base delay and random delay. For example, if the base delay is 3 and <b>pcmk_delay_max</b> is 10, the random delay will be between 3 and 10. For general information about fencing delay parameters and their interactions, see <a href="#">Fencing delays</a> .
<b>pcmk_action_limit</b>	integer	1	The maximum number of actions that can be performed in parallel on this device. The cluster property <b>concurrent-fencing=true</b> needs to be configured first (this is the default value). A value of -1 is unlimited.
<b>pcmk_on_action</b>	string	on	For advanced use only: An alternate command to run instead of <b>on</b> . Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific, command that implements the <b>on</b> action.
<b>pcmk_on_timeout</b>	time	60s	For advanced use only: Specify an alternate timeout to use for <b>on</b> actions instead of <b>stonith-timeout</b> . Some devices need much more or much less time to complete than normal. Use this to specify an alternate, device-specific, timeout for <b>on</b> actions.
<b>pcmk_on_retries</b>	integer	2	For advanced use only: The maximum number of times to retry the <b>on</b> command within the timeout period. Some devices do not support multiple connections. Operations may <b>fail</b> if the device is busy with another task so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries <b>on</b> actions before giving up.

In addition to the properties you can set for individual fence devices, there are also cluster properties you can set that determine fencing behavior, as described in the following table.

Table 4.4. Cluster Properties that Determine Fencing Behavior

Option	Default	Description

Option	Default	Description
<b>stonith-enabled</b>	true	<p>Indicates that failed nodes and nodes with resources that cannot be stopped should be fenced. Protecting your data requires that you set this <b>true</b>.</p> <p>If <b>true</b>, or unset, the cluster will refuse to start resources unless one or more STONITH resources have been configured also.</p> <p>Red Hat only supports clusters with this value set to <b>true</b>.</p>
<b>stonith-action</b>	reboot	Action to send to fencing device. Allowed values: <b>reboot, off</b> . The value <b>poweroff</b> is also allowed, but is only used for legacy devices.
<b>stonith-timeout</b>	60s	How long to wait for a STONITH action to complete.
<b>stonith-max-attempts</b>	10	How many times fencing can fail for a target before the cluster will no longer immediately re-attempt it.
<b>stonith-watchdog-timeout</b>		The maximum time to wait until a node can be assumed to have been killed by the hardware watchdog. It is recommended that this value be set to twice the value of the hardware watchdog timeout. This option is needed only if watchdog-only SBD configuration is used for fencing.
<b>concurrent-fencing</b>	true	Allow fencing operations to be performed in parallel.

Option	Default	Description
<b>fence-reaction</b>	stop	<p>Determines how a cluster node should react if notified of its own fencing. A cluster node may receive notification of its own fencing if fencing is misconfigured, or if fabric fencing is in use that does not cut cluster communication. Allowed values are <b>stop</b> to attempt to immediately stop Pacemaker and stay stopped, or <b>panic</b> to attempt to immediately reboot the local node, falling back to stop on failure.</p> <p>Although the default value for this property is <b>stop</b>, the safest choice for this value is <b>panic</b>, which attempts to immediately reboot the local node. If you prefer the stop behavior, as is most likely to be the case in conjunction with fabric fencing, it is recommended that you set this explicitly.</p>
<b>priority-fencing-delay</b>	0 (disabled)	<p>Sets a fencing delay that allows you to configure a two-node cluster so that in a split-brain situation the node with the fewest or least important resources running is the node that gets fenced. For general information about fencing delay parameters and their interactions, see <a href="#">Fencing delays</a>.</p>

For information about setting cluster properties, see

[https://docs.redhat.com/en/documentation/red\\_hat\\_enterprise\\_linux/10/html/configuring\\_and\\_managing\\_cluster-properties](https://docs.redhat.com/en/documentation/red_hat_enterprise_linux/10/html/configuring_and_managing_cluster-properties)

#### 4.9.4. Fencing delays

In two-node clusters, simultaneous communication loss can cause nodes to fence each other, shutting down the entire cluster. Configure a fencing delay to prevent this race condition. Delays are unnecessary in larger clusters, where quorum determines fencing authority.

You can set different types of fencing delays, depending on your system requirements.

- **static fencing delays**

A static fencing delay is a fixed, predetermined delay. Setting a static delay on one node makes that node more likely to be fenced because it increases the chances that the other node will initiate fencing first after detecting lost communication. In an active/passive cluster, setting a delay on a passive node makes it more likely that the passive node will be fenced when communication breaks down. You configure a static delay by using the **pcs\_delay\_base** cluster property. You can set this property when a separate fence device is used for each node or when a single fence device is used for all nodes.

- **dynamic fencing delays**

A dynamic fencing delay is random. It can vary and is determined at the time fencing is needed.

You configure a random delay and specify a maximum value for the combined base delay and random delay with the **pcs\_delay\_max** cluster property. When the fencing delay for each node is random, which node is fenced is also random. You may find this feature useful if your cluster is configured with a single fence device for all nodes in an active/active design.

- **priority fencing delays**

A priority fencing delay is based on active resource priorities. If all resources have the same priority, the node with the fewest resources running is the node that gets fenced. In most cases, you use only one delay-related parameter, but it is possible to combine them. Combining delay-related parameters adds the priority values for the resources together to create a total delay. You configure a priority fencing delay with the **priority-fencing-delay** cluster property. You may find this feature useful in an active/active cluster design because it can make the node running the fewest resources more likely to be fenced when communication between the nodes is lost.

### The **pcmk\_delay\_base** cluster property

Setting the **pcmk\_delay\_base** cluster property enables a base delay for fencing and specifies a base delay value.

When you set the **pcmk\_delay\_max** cluster property in addition to the **pcmk\_delay\_base** property, the overall delay is derived from a random delay value added to this static delay so that the sum is kept below the maximum delay. When you set **pcmk\_delay\_base** but do not set **pcmk\_delay\_max**, there is no random component to the delay and it will be the value of **pcmk\_delay\_base**.

You can specify different values for different nodes with the **pcmk\_delay\_base** parameter. This allows a single fence device to be used in a two-node cluster, with a different delay for each node. You do not need to configure two separate devices to use separate delays. To specify different values for different nodes, you map the host names to the delay value for that node using a similar syntax to **pcmk\_host\_map**. For example, **node1:0;node2:10s** would use no delay when fencing **node1** and a 10-second delay when fencing **node2**.

### The **pcmk\_delay\_max** cluster property

Setting the **pcmk\_delay\_max** cluster property enables a random delay for fencing actions and specifies the maximum delay, which is the maximum value of the combined base delay and random delay. For example, if the base delay is 3 and **pcmk\_delay\_max** is 10, the random delay will be between 3 and 10.

When you set the **pcmk\_delay\_base** cluster property in addition to the **pcmk\_delay\_max** property, the overall delay is derived from a random delay value added to this static delay so that the sum is kept below the maximum delay. When you set **pcmk\_delay\_max** but do not set **pcmk\_delay\_base** there is no static component to the delay.

### The **priority-fencing-delay** cluster property

Setting the **priority-fencing-delay** cluster property allows you to configure a two-node cluster so that in a split-brain situation the node with the fewest or least important resources running is the node that gets fenced.

The **priority-fencing-delay** property can be set to a time duration. The default value for this property is 0 (disabled). If this property is set to a non-zero value, and the priority meta-attribute is configured for at least one resource, then in a split-brain situation the node with the highest combined priority of all resources running on it will be more likely to remain operational. For example, if you set **pcs resource defaults update priority=1** and **pcs property set priority-fencing-delay=15s** and no other priorities

are set, then the node running the most resources will be more likely to remain operational because the other node will wait 15 seconds before initiating fencing. If a particular resource is more important than the rest, you can give it a higher priority.

The node running the promoted role of a promotable clone gets an extra 1 point if a priority has been configured for that clone.

### Interaction of fencing delays

Setting more than one type of fencing delay yields the following results:

- Any delay set with the **priority-fencing-delay** property is added to any delay from the **pcmk\_delay\_base** and **pcmk\_delay\_max** fence device properties. This behavior allows some delay when both nodes have equal priority, or both nodes need to be fenced for some reason other than node loss, as when **on-fail=fencing** is set for a resource monitor operation. When setting these delays in combination, set the **priority-fencing-delay** property to a value that is significantly greater than the maximum delay from **pcmk\_delay\_base** and **pcmk\_delay\_max** to be sure the prioritized node is preferred. Setting this property to twice this value is always safe.
- Only fencing scheduled by Pacemaker itself observes fencing delays. Fencing scheduled by external code such as **dlm\_controlld** and fencing implemented by the **pcs stonith fence** command do not provide the necessary information to the fence device.
- Some individual fence agents implement a delay parameter, with a name determined by the agent and which is independent of delays configured with a **pcmk\_delay\_\*** property. If both of these delays are configured, they are added together and would generally not be used in conjunction.

#### 4.9.5. Testing a fence device

Validate fence devices to ensure the cluster can successfully recover from node failures and prevent data corruption. A complete testing strategy involves verifying network connectivity, executing the fence agent script directly, triggering the fence action through the cluster manager, and simulating a physical node failure.



#### NOTE

When a Pacemaker cluster node or Pacemaker remote node is fenced a hard kill should occur and not a graceful shutdown of the operating system. If a graceful shutdown occurs when your system fences a node, disable ACPI soft-off in the **/etc/systemd/logind.conf** file so that your system ignores any power-button-pressed signal. For instructions on disabling ACPI soft-off in the **logind.conf** file, see [Disabling ACPI soft-off in the logind.conf file](#)

Use the following procedure to test a fence device.

#### Procedure

1. Use SSH, Telnet, HTTP, or whatever remote protocol is used to connect to the device to manually log in and test the fence device or see what output is given. For example, if you will be configuring fencing for an IPMI-enabled device, then try to log in remotely with **ipmitool**. Take note of the options used when logging in manually because those options might be needed when using the fencing agent.

If you are unable to log in to the fence device, verify that the device is pingable, there is nothing such as a firewall configuration that is preventing access to the fence device, remote access is enabled on the fencing device, and the credentials are correct.

2. Run the fence agent manually, using the fence agent script. This does not require that the cluster services are running, so you can perform this step before the device is configured in the cluster. This can ensure that the fence device is responding properly before proceeding.



### NOTE

These examples use the **fence\_ipmilan** fence agent script for an iLO device. The actual fence agent you will use and the command that calls that agent will depend on your server hardware. You should consult the man page for the fence agent you are using to determine which options to specify. You will usually need to know the login and password for the fence device and other information related to the fence device.

The following example shows the format you would use to run the **fence\_ipmilan** fence agent script with **-o status** parameter to check the status of the fence device interface on another node without actually fencing it. This allows you to test the device and get it working before attempting to reboot the node. When running this command, you specify the name and password of an iLO user that has power on and off permissions for the iLO device.

```
# fence_ipmilan -a ipaddress -l username -p password -o status
```

The following example shows the format you would use to run the **fence\_ipmilan** fence agent script with the **-o reboot** parameter. Running this command on one node reboots the node managed by this iLO device.

```
# fence_ipmilan -a ipaddress -l username -p password -o reboot
```

If the fence agent failed to properly do a status, off, on, or reboot action, you should check the hardware, the configuration of the fence device, and the syntax of your commands. In addition, you can run the fence agent script with the debug output enabled. The debug output is useful for some fencing agents to see where in the sequence of events the fencing agent script is failing when logging into the fence device.

```
# fence_ipmilan -a ipaddress -l username -p password -o status -D /tmp/${hostname}-fence_agent.debug
```

When diagnosing a failure that has occurred, you should ensure that the options you specified when manually logging in to the fence device are identical to what you passed on to the fence agent with the fence agent script.

For fence agents that support an encrypted connection, you may see an error due to certificate validation failing, requiring that you trust the host or that you use the fence agent's **ssl-insecure** parameter. Similarly, if SSL/TLS is disabled on the target device, you may need to account for this when setting the SSL parameters for the fence agent.



## NOTE

If the fence agent that is being tested is a **fence\_drac**, **fence\_ilc**, or some other fencing agent for a systems management device that continues to fail, then fall back to trying **fence\_ipmilan**. Most systems management cards support IPMI remote login and the only supported fencing agent is **fence\_ipmilan**.

3. Once the fence device has been configured in the cluster with the same options that worked manually and the cluster has been started, test fencing with the **pcs stonith fence** command from any node (or even multiple times from different nodes), as in the following example. The **pcs stonith fence** command reads the cluster configuration from the CIB and calls the fence agent as configured to execute the fence action. This verifies that the cluster configuration is correct.

```
# pcs stonith fence node_name
```

If the **pcs stonith fence** command works properly, that means the fencing configuration for the cluster should work when a fence event occurs. If the command fails, it means that cluster management cannot invoke the fence device through the configuration it has retrieved. Check for the following issues and update your cluster configuration as needed.

- Check your fence configuration. For example, if you have used a host map you should ensure that the system can find the node using the host name you have provided.
- Check whether the password and user name for the device include any special characters that could be misinterpreted by the bash shell. Making sure that you enter passwords and user names surrounded by quotation marks could address this issue.
- Check whether you can connect to the device using the exact IP address or host name you specified in the **pcs stonith** command. For example, if you give the host name in the stonith command but test by using the IP address, that is not a valid test.
- If the protocol that your fence device uses is accessible to you, use that protocol to try to connect to the device. For example many agents use ssh or telnet. You should try to connect to the device with the credentials you provided when configuring the device, to see if you get a valid prompt and can log in to the device.

If you determine that all your parameters are appropriate but you still have trouble connecting to your fence device, you can check the logging on the fence device itself, if the device provides that, which will show if the user has connected and what command the user issued. You can also search through the **/var/log/messages** file for instances of stonith and error, which could give some idea of what is transpiring, but some agents can provide additional information.

4. Once the fence device tests are working and the cluster is up and running, test an actual failure. To do this, take an action in the cluster that should initiate a token loss.
  - Take down a network. How you take a network depends on your specific configuration. In many cases, you can physically pull the network or power cables out of the host. For information about simulating a network failure, see the Red Hat Knowledgebase solution [What is the proper way to simulate a network failure on a RHEL Cluster?](#) .

**NOTE**

Disabling the network interface on the local host rather than physically disconnecting the network or power cables is not recommended as a test of fencing because it does not accurately simulate a typical real-world failure.

- Block corosync traffic both inbound and outbound using the local firewall. The following example blocks corosync, assuming the default corosync port is used, **firewalld** is used as the local firewall, and the network interface used by corosync is in the default firewall zone:

```
# firewall-cmd --direct --add-rule ipv4 filter OUTPUT 2 -p udp --dport=5405 -j DROP
# firewall-cmd --add-rich-rule='rule family="ipv4" port port="5405" protocol="udp" drop'
```

- Simulate a crash and panic your machine with **sysrq-trigger**. Note, however, that triggering a kernel panic can cause data loss; it is recommended that you disable your cluster resources first.

```
# echo c > /proc/sysrq-trigger
```

#### 4.9.6. Configuring fencing levels

Pacemaker supports fencing nodes with multiple devices through a feature called fencing topologies. To implement topologies, create the individual devices as you normally would and then define one or more fencing levels in the fencing topology section in the configuration.

Pacemaker processes fencing levels as follows:

- Each level is attempted in ascending numeric order, starting at 1.
- If a device fails, processing terminates for the current level. No further devices in that level are exercised and the next level is attempted instead.
- If all devices are successfully fenced, then that level has succeeded and no other levels are tried.
- The operation is finished when a level has passed (success), or all levels have been attempted (failed).

Use the following command to add a fencing level to a node. The devices are given as a comma-separated list of **stonith** ids, which are attempted for the node at that level.

```
pcs stonith level add level node devices
```

The following example sets up fence levels so that if the device **my\_il0** fails and is unable to fence the node, then Pacemaker attempts to use the device **my\_apc**.

#### Prerequisites

- You have configured an ilo fence device called **my\_il0** for node **rh7-2**.
- You have configured an apc fence device called **my\_apc** for node **rh7-2**.

## Procedure

1. Add a fencing level of 1 for fence device **my\_il0** on node **rh7-2**.

```
# pcs stonith level add 1 rh7-2 my_il0
```

2. Add a fencing level of 2 for fence device **my\_apc** on node **rh7-2**.

```
# pcs stonith level add 2 rh7-2 my_apc
```

3. List the currently configured fencing levels.

```
# pcs stonith level
```

Node: rh7-2

  Level 1 - my\_il0

  Level 2 - my\_apc

## Additional resources

[Configuring node-specific values using node attributes](#) .

### 4.9.7. Removing a fence level

You can remove the fence level for the specified node and device. If no nodes or devices are specified then the fence level you specify is removed from all nodes.

## Procedure

- Remove the fence level for the specified node and device:

```
# pcs stonith level remove level [node_id] [stonith_id] ... [stonith_id]
```

### 4.9.8. Clearing fence levels

You can clear the fence levels on the specified node or stonith id. If you do not specify a node or stonith id, all fence levels are cleared.

## Procedure

- Clear the fence levels on the specified node or stonith id:

```
# pcs stonith level clear [node]|stonith_id(s)]
```

- If you specify more than one stonith id, they must be separated by a comma and no spaces, as in the following example.

```
# pcs stonith level clear dev_a,dev_b
```

### 4.9.9. Verifying nodes and devices in fence levels

You can verify that all fence devices and nodes specified in fence levels exist.

## Procedure

- Use the following command to verify that all fence devices and nodes specified in fence levels exist:

```
# pcs stonith level verify
```

### 4.9.10. Specifying nodes in fencing topology

You can specify nodes in fencing topology by a regular expression applied on a node name and by a node attribute and its value.

## Procedure

- The following commands configure nodes **node1**, **node2**, and **node3** to use fence devices **apc1** and **apc2**, and nodes **node4**, **node5**, and **node6** to use fence devices **apc3** and **apc4**:

```
# pcs stonith level add 1 "regexp%node[1-3]" apc1,apc2
# pcs stonith level add 1 "regexp%node[4-6]" apc3,apc4
```

- The following commands yield the same results by using node attribute matching:

```
# pcs node attribute node1 rack=1
# pcs node attribute node2 rack=1
# pcs node attribute node3 rack=1
# pcs node attribute node4 rack=2
# pcs node attribute node5 rack=2
# pcs node attribute node6 rack=2
# pcs stonith level add 1 attrib%rack=1 apc1,apc2
# pcs stonith level add 1 attrib%rack=2 apc3,apc4
```

### 4.9.11. Configuring fencing for redundant power supplies

When configuring fencing for redundant power supplies, the cluster must ensure that when attempting to reboot a host, both power supplies are turned off before either power supply is turned back on.

If the node never completely loses power, the node may not release its resources. This opens up the possibility of nodes accessing these resources simultaneously and corrupting them.

You need to define each device only once and to specify that both are required to fence the node.

## Procedure

1. Create the first fence device.

```
# pcs stonith create apc1 fence_apc_snmp ipaddr=apc1.example.com login=user
passwd='7a4D#1j!pz864'
pcmk_host_map="node1.example.com:1;node2.example.com:2"
```

2. Create the second fence device.

```
# pcs stonith create apc2 fence_apc_snmp ipaddr=apc2.example.com login=user
passwd='7a4D#1j!pz864'
pcmk_host_map="node1.example.com:1;node2.example.com:2"
```

3. Specify that both devices are required to fence the node.

```
# pcs stonith level add 1 node1.example.com apc1,apc2
# pcs stonith level add 1 node2.example.com apc1,apc2
```

## 4.9.12. Administering fence devices

The **pcs** command-line interface provides a variety of commands you can use to administer your fence devices after you have configured them.

### 4.9.12.1. Displaying configured fence devices

The following command shows all currently configured fence devices. If a *stonith\_id* is specified, the command shows the options for that configured fencing device only. If the **--full** option is specified, all configured fencing options are displayed.

```
pcs stonith config [stonith_id] [--full]
```

### 4.9.12.2. Exporting fence devices as pcs commands

You can display the **pcs** commands that can be used to re-create configured fence devices on a different system using the **--output-format=cmd** option of the **pcs stonith config** command.

The following commands create a **fence\_apc\_snmp** fence device and display the **pcs** command you can use to re-create the device.

```
# pcs stonith create myapc fence_apc_snmp ip="zapc.example.com"
pcmk_host_map="z1.example.com:1;z2.example.com:2" username="apc" password="apc"
# pcs stonith config --output-format=cmd
Warning: Only 'text' output format is supported for stonith levels
pcs stonith create --no-default-ops --force -- myapc fence_apc_snmp \
    ip=zapc.example.com password=apc 'pcmk_host_map=z1.example.com:1;z2.example.com:2'
    username=apc \
    op \
        monitor interval=60s id=myapc-monitor-interval-60s
```

### 4.9.12.3. Exporting fence level configuration

The **pcs stonith config** and the **pcs stonith level config** commands support the **--output-format=** option to export the fencing level configuration in JSON format and as **pcs** commands.

- Specifying **--output-format=cmd** displays the **pcs** commands created from the current cluster configuration that configure fencing levels. You can use these commands to re-create configured fencing levels on a different system.
- Specifying **--output-format=json** displays the fencing level configuration in JSON format, which is suitable for machine parsing.

#### 4.9.12.4. Modifying and deleting fence devices

Modify or add options to a currently configured fencing device with the following command.

```
pcs stonith update stonith_id [stonith_device_options]
```

Updating a SCSI fencing device with the **pcs stonith update** command causes a restart of all resources running on the same node where the fencing resource was running. You can use either version of the following command to update SCSI devices without causing a restart of other cluster resources. SCSI fencing devices can be configured as multipath devices.

```
pcs stonith update-scsi-devices stonith_id set device-path1 device-path2  
pcs stonith update-scsi-devices stonith_id add device-path1 remove device-path2
```

Use the following command to remove a fencing device from the current configuration.

```
pcs stonith delete stonith_id
```

#### 4.9.12.5. Manually fencing a cluster node

You can fence a node manually with the following command. If you specify the **--off** option this will use the **off** API call to stonith which will turn the node off instead of rebooting it.

```
pcs stonith fence node [--off]
```

In a situation where no fence device is able to fence a node even if it is no longer active, the cluster may not be able to recover the resources on the node. If this occurs, after manually ensuring that the node is powered down you can enter the following command to confirm to the cluster that the node is powered down and free its resources for recovery.



#### WARNING

If the node you specify is not actually off, but running the cluster software or services normally controlled by the cluster, data corruption and cluster failure occurs.

```
pcs stonith confirm node
```

#### 4.9.12.6. Disabling a fence device

To disable a fencing device, run the **pcs stonith disable** command.

The following command disables the fence device **myapc**.

```
# pcs stonith disable myapc
```

#### 4.9.12.7. Preventing a node from using a fencing device

To prevent a specific node from using a fencing device, you can configure location constraints for the fencing resource.

The following example prevents fence device **node1-ipmi** from running on **node1**.

```
# pcs constraint location node1-ipmi avoids node1
```

#### 4.9.13. Configuring ACPI for use with integrated fence devices

If your cluster uses integrated fence devices, you must configure ACPI (Advanced Configuration and Power Interface) to ensure immediate and complete fencing.

If a cluster node is configured to be fenced by an integrated fence device, disable ACPI Soft-Off for that node. Disabling ACPI Soft-Off allows an integrated fence device to turn off a node immediately and completely rather than attempting a clean shutdown (for example, **shutdown -h now**). Otherwise, if ACPI Soft-Off is enabled, an integrated fence device can take four or more seconds to turn off a node (see the note that follows). In addition, if ACPI Soft-Off is enabled and a node panics or freezes during shutdown, an integrated fence device may not be able to turn off the node. Under those circumstances, fencing is delayed or unsuccessful. Consequently, when a node is fenced with an integrated fence device and ACPI Soft-Off is enabled, a cluster recovers slowly or requires administrative intervention to recover.



##### NOTE

The amount of time required to fence a node depends on the integrated fence device used. Some integrated fence devices perform the equivalent of pressing and holding the power button; therefore, the fence device turns off the node in four to five seconds. Other integrated fence devices perform the equivalent of pressing the power button momentarily, relying on the operating system to turn off the node; therefore, the fence device turns off the node in a time span much longer than four to five seconds.

- The preferred way to disable ACPI Soft-Off is to change the BIOS setting to "instant-off" or an equivalent setting that turns off the node without delay, as described in [Disabling ACPI Soft-Off with the Bios](#).

Disabling ACPI Soft-Off with the BIOS may not be possible with some systems. If disabling ACPI Soft-Off with the BIOS is not satisfactory for your cluster, you can disable ACPI Soft-Off with one of the following alternate methods:

- Setting **HandlePowerKey=ignore** in the **/etc/systemd/logind.conf** file and verifying that the node node turns off immediately when fenced, as described in [Disabling ACPI soft-off in the logind.conf file](#). This is the first alternate method of disabling ACPI Soft-Off.
- Appending **acpi=off** to the kernel boot command line, as described in [Disabling ACPI completely in the GRUB 2 file](#). This is the second alternate method of disabling ACPI Soft-Off, if the preferred or the first alternate method is not available.



##### IMPORTANT

This method completely disables ACPI; some computers do not boot correctly if ACPI is completely disabled. Use this method *only* if the other methods are not effective for your cluster.

#### 4.9.13.1. Disabling ACPI Soft-Off with the BIOS

You can disable ACPI Soft-Off by configuring the BIOS of each cluster node.



##### NOTE

The procedure for disabling ACPI Soft-Off with the BIOS may differ among server systems. You should verify this procedure with your hardware documentation.

##### Procedure

1. Reboot the node and start the **BIOS CMOS Setup Utility** program.
2. Navigate to the Power menu (or equivalent power management menu).
3. At the Power menu, set the **Soft-Off by PWR-BTTN** function (or equivalent) to **Instant-Off** (or the equivalent setting that turns off the node by means of the power button without delay). The **BIOS CMOS Setup Utiliy** example below shows a Power menu with **ACPI Function** set to **Enabled** and **Soft-Off by PWR-BTTN** set to **Instant-Off**.



##### NOTE

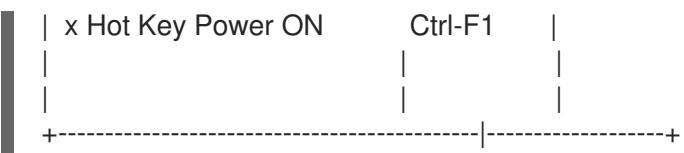
The equivalents to **ACPI Function**, **Soft-Off by PWR-BTTN**, and **Instant-Off** may vary among computers. However, the objective of this procedure is to configure the BIOS so that the computer is turned off by means of the power button without delay.

4. Exit the **BIOS CMOS Setup Utility** program, saving the BIOS configuration.
5. Verify that the node turns off immediately when fenced. For information about testing a fence device, see [Testing a fence device](#).

##### BIOS CMOS Setup Utility:

`Soft-Off by PWR-BTTN` set to  
`Instant-Off`

ACPI Function	[Enabled]		Item Help	
ACPI Suspend Type	[S1(POS)]	-----		
x Run VGABIOS if S3 Resume	Auto		Menu Level	*
Suspend Mode	[Disabled]			
HDD Power Down	[Disabled]			
Soft-Off by PWR-BTTN	[Instant-Off]			
CPU THR-M-Throttling	[50.0%]			
Wake-Up by PCI card	[Enabled]			
Power On by Ring	[Enabled]			
Wake Up On LAN	[Enabled]			
x USB KB Wake-Up From S3	Disabled			
Resume by Alarm	[Disabled]			
x Date(of Month) Alarm	0			
x Time(hh:mm:ss) Alarm	0 : 0 :			
POWER ON Function	[BUTTON ONLY]			
x KB Power ON Password	Enter			



This example shows **ACPI Function** set to **Enabled**, and **Soft-Off by PWR-BTTN** set to **Instant-Off**.

#### 4.9.13.2. Disabling ACPI Soft-Off in the logind.conf file

You can disable power-key handing in the **/etc/systemd/logind.conf** file.

##### Procedure

1. Define the following configuration in the **/etc/systemd/logind.conf** file:

```
HandlePowerKey=ignore
```

2. Restart the **systemd-logind** service:

```
# systemctl restart systemd-logind.service
```

##### Verification

1. Verify that the node turns off immediately when fenced. For information about testing a fence device, see [Testing a fence device](#).

#### 4.9.13.3. Disabling ACPI completely in the GRUB 2 file

You can disable ACPI Soft-Off by appending **acpi=off** to the GRUB menu entry for a kernel.



##### IMPORTANT

This method completely disables ACPI; some computers do not boot correctly if ACPI is completely disabled. Use this method *only* if the other methods are not effective for your cluster.

##### Procedure

1. Use the **--args** option in combination with the **--update-kernel** option of the **grubby** tool to change the **grub.cfg** file of each cluster node as follows:

```
# grubby --args=acpi=off --update-kernel=ALL
```

2. Reboot the node.

##### Verification

1. Verify that the node turns off immediately when fenced. For information about testing a fence device, see [Testing a fence device](#)

## 4.10. CONFIGURING THE VIRTUAL IP MANAGEMENT RESOURCE AGENT

The **gcp-vpc-move-vip** resource agent attaches a secondary IP address (alias IP) to a running instance. You can assign this floating IP address between different nodes in the cluster.

```
# pcs resource describe gcp-vpc-move-vip
```

You can configure the resource agent to use a primary subnet address range or a secondary subnet address range.

### 4.10.1. Configuring the primary subnet address range

If you need to automate or manage the assigned IP addresses allocation for VM or other resources within a subnet, you can use the primary subnet address range. It ensures that the primary address range is correctly set and configure to use as stable IP addresses for the primary virtual private network (VPC) subnet.

#### Procedure

1. Create the **aliasip** resource by including an unused internal IP address and the CIDR block:

```
# pcs resource create aliasip gcp-vpc-move-vip alias_ip=UnusedIPaddress/CIDRblock
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource create aliasip gcp-vpc-move-vip
alias_ip=10.10.10.200/32
```

2. Create an **IPAddr2** resource for managing the IP on the node:

```
# pcs resource create vip IPAddr2 nic=interface ip=AliasIPaddress cidr_netmask=32
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource create vip IPAddr2 nic=eth0 ip=10.10.10.200
cidr_netmask=32
```

3. Group the network resources under **vipgrp**:

```
# pcs resource group add vipgrp aliasip vip
```

#### Verification

1. Verify the active resources and under the **vipgrp** group:

```
# pcs status
```

2. Verify the movable resources across the nodes:

```
# pcs resource move vip Node
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource move vip rhel81-node-03
```

3. Verify if the **vip** successfully started on a different node:

```
# pcs status
```

#### 4.10.2. Configuring the secondary subnet address range

You can use the secondary subnet address range if you need to assign IP addresses from additional and predefined range within the same subnet, without creating a new subnet. It is useful for specific purposes such as custom routing. With secondary subnet address range, you can manage network traffic in a single subnet with multiple IP address ranges.

##### Prerequisites

- You have created a custom network and a subnet
- Optional: You have installed Google Cloud SDK. For instructions, see [Installing the Google Cloud SDK](#).



##### NOTE

You can also use the **gcloud** commands in the following procedure in the terminal that you can activate in the Google Cloud web console.

##### Procedure

1. Create a secondary subnet address range:

```
# gcloud compute networks subnets update SubnetName --region RegionName --add-secondary-ranges SecondarySubnetName=SecondarySubnetRange
```

Example:

```
# gcloud compute networks subnets update range0 --region us-west1 --add-secondary-ranges range1=10.10.20.0/24
```

2. Create the **aliasip** resource with an unused internal IP address in the secondary subnet address range and the CIDR block:

```
# pcs resource create aliasip gcp-vpc-move-vip alias_ip=UnusedIPaddress/CIDRblock
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource create aliasip gcp-vpc-move-vip alias_ip=10.10.20.200/32
```

3. Create an **IPAddr2** resource for managing the IP on the node:

```
# pcs resource create vip IPAddr2 nic=interface ip=AliasIPaddress cidr_netmask=32
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource create vip IPaddr2 nic=eth0 ip=10.10.20.200  
cidr_netmask=32
```

4. Group the network resources under **vipgrp**:

```
# pcs resource group add vipgrp aliasip vip
```

## Verification

1. Verify that the active resources are under the **vipgrp** group:

```
# pcs status
```

2. Verify that the movable resources across different nodes:

```
# pcs resource move vip Node
```

Example:

```
[root@rhel81-node-01 ~]# pcs resource move vip rhel81-node-03
```

3. Verify that the **vip** successfully started on a different node:

```
# pcs status
```

### 4.10.3. Additional resources

- [Support Policies for Red Hat High Availability clusters - Transport Protocols](#)
- [Exploring Red Hat High Availability's Components, Concepts, and Features - Overview of Transport Protocols](#)
- [Design Guidance for Red Hat High Availability Clusters - Selecting the Transport Protocol](#)

# CHAPTER 5. CONFIGURING RHEL ON GOOGLE CLOUD WITH SECURE BOOT

To provide security at boot time for Red Hat Enterprise Linux (RHEL) instances on Google Cloud, you can configure Secure Boot. During startup, Secure Boot verifies digital signatures of the boot loader and other components. Then, Secure Boot allows only trusted programs to load and blocks unauthorized programs.

You can set up RHEL instances by using either Google Cloud Marketplace images or custom RHEL images with Secure Boot enabled.

## 5.1. TYPES OF RHEL IMAGES ON GOOGLE CLOUD

### Google Cloud Marketplace images

The Google Cloud Marketplace provides a pre-configured Red Hat Enterprise Linux (RHEL) image that is tailored for specified use cases such as data processing, system management, and web development. This type of ready-to-use image reduces setup by minimizing manual installation and configuration time required for operating systems and software packages.

### Custom RHEL image

A custom RHEL image offers flexibility to customers and organizations to build and deploy tailored environments that meet specific application and workflow requirements. By creating a custom RHEL image, you can use RHEL instances that are pre-installed with necessary tools, configurations, and security policies. This customization aims at greater control over the infrastructure.

## 5.2. UNDERSTANDING SECURE BOOT FOR RHEL ON CLOUD

Secure Boot is a Unified Extensible Firmware Interface (UEFI) feature that verifies digital signatures of boot components, such as boot loader and kernel, against trusted keys stored in hardware. Secure Boot prevents unauthorized or tampered software from running during boot, protecting your system from malicious code.

If Secure Boot detects any tampered components or components signed by untrusted entities, it aborts the boot process. Secure Boot plays a critical role in configuring a Confidential Virtual Machine (CVM) by ensuring that only trusted entities participate in the boot chain. It authenticates access to specific device paths through defined interfaces, enforces the use of the latest configuration, and permanently overwrites earlier configurations. When the Red Hat Enterprise Linux (RHEL) kernel boots with Secure Boot enabled, it enters the **lockdown** mode, allowing only kernel modules signed by a trusted vendor to load. Therefore, Secure Boot strengthens the security of the operating system boot sequence.

### 5.2.1. Components of Secure Boot

The Secure Boot mechanism consists of firmware, signature databases, cryptographic keys, boot loader, hardware modules, and the operating system. The following are the components of the UEFI trusted variables:

- Key Exchange Key database (KEK): An exchange of public keys to establish trust between the RHEL operating system and the VM firmware. You can also update Allowed Signature database (**db**) and Forbidden Signature database (**dbx**) by using these keys.
- Platform Key database (PK): A self-signed single-key database to establish trust between the VM firmware and the cloud platform. The PK also updates the KEK database.
- Allowed Signature database (**db**): A database that maintains a list of certificates or binary

hashes to check whether the binary file can boot on the system. Additionally, all certificates from **db** are imported to the **.platform** keyring of the RHEL kernel. With this feature, you can add and load signed third party kernel modules in the **lockdown** mode.

- **Forbidden Signature database (dbx):** A database that maintains a list of certificates or binary hashes that are not allowed to boot on the system.



#### NOTE

Binary files check against the **dbx** database and the Secure Boot Advanced Targeting (SBAT) mechanism. With SBAT, you can revoke older versions of specific binaries by keeping the certificate that has signed binaries as valid.

### 5.2.2. Stages of Secure Boot for RHEL on Cloud

When a RHEL instance boots in the Unified Kernel Image (UKI) mode and with Secure Boot enabled, the RHEL instance interacts with the cloud service infrastructure in the following sequence:

1. *Initialization:* When a RHEL instance boots, the cloud-hosted firmware initially boots and implements the Secure Boot mechanism.
2. *Variable store initialization:* The firmware initializes UEFI variables from a variable store, a dedicated storage area for information that firmware needs to manage for the boot process and runtime operations. When the RHEL instance boots for the first time, the store initializes from default values associated with the VM image.
3. *Boot loader:* When booted, the firmware loads the first stage boot loader. For the RHEL instance in a x86 UEFI environment, the first stage boot loader is shim. The shim boot loader authenticates and loads the next stage of the boot process and acts as a bridge between UEFI and GRUB.
  - a. The shim x86 binary in RHEL is currently signed by the **Microsoft Corporation UEFI CA 2011** Microsoft certificate so that the RHEL instance can boot in the Secure Boot enabled mode on various hardware and virtualized platforms where the Allowed Signature database (**db**) has the default Microsoft certificates.
  - b. The shim binary extends the list of trusted certificates with Red Hat Secure Boot CA and optionally, with Machine Owner Key (**MOK**).
4. *UKI:* The shim binary loads the RHEL UKI (the **kernel-uki-virt** package). The corresponding certificate, **Red Hat Secure Boot Signing 504** on the x86\_64 architecture, signs the UKI. You can find this certificate in the **redhat-sb-certs** package. Red Hat Secure Boot CA signs this certificate, so the check succeeds.
5. *UKI add-ons:* When you use the UKI **cmdline** extensions, the RHEL kernel actively checks their signatures against **db**, **MOK**, and certificates shipped with shim. This process ensures that either the operating system vendor RHEL or a user has signed the extensions.

When the RHEL kernel boots in the Secure Boot mode, it enters **lockdown** mode. After entering **lockdown**, the RHEL kernel adds the **db** keys to the **.platform** keyring and the **MOK** keys to the **.machine** keyring. During the kernel build process, the build system works with an ephemeral key, which consists of private and public keys. The build system signs standard RHEL kernel modules, such as **kernel-modules-core**, **kernel-modules**, and **kernel-modules-extra**. After the completion of each kernel build, the private key becomes obsolete to sign third-party modules. You can use certificates from **db** and **MOK** for this purpose.

## Additional resources

- [SBAT mechanism](#)
- [Red Hat Enterprise Linux and Secure Boot in cloud](#)

## 5.3. CONFIGURING A RHEL INSTANCE FROM A CUSTOM RHEL IMAGE WITH SECURE BOOT

To ensure that your Red Hat Enterprise Linux (RHEL) instance on Google Cloud has a secured operating system booting process, use the Secure Boot mechanism. You can configure the Secure Boot during registration of a custom RHEL Google Cloud image.

This image consists of pre-stored Unified Extensible Firmware Interface (UEFI) variables. Therefore, instances launched from this image use the Secure Boot mechanism with the required variables during the first boot. You can also include a custom certificate in the SecureBoot DB to sign custom artifacts, such as third party kernel modules and Unified Kernel Image (UKI) add-ons.

### Prerequisites

1. You have created and uploaded a RHEL Google Cloud image. For details, see [Uploading images to Google Cloud](#).
2. You have installed the following packages:
  - **python3**
  - **efivar**
  - **keyutils**
  - **openssl**
  - **python3-virt-firmware**
3. You have installed the **google-cloud-cli** utility. For details, see [installing gcloud CLI on RHEL](#).

### Procedure

1. Create a new random Universally Unique Identifier (UUID) and store it in a system generated random text file:

```
$ uuidgen --random > GUID.txt
```

2. Generate a new RSA private key **PK.key** and a self-signed X.509 certificate **PK.cer** for the Platform Key database:

```
$ openssl req -quiet \  
-newkey rsa:4096 \  
-nodes -keyout PK.key \  
-new -x509 -sha256 \  
-days 3650 \  
-subj "/CN=Platform key/" \  
-outform DER \  
-out PK.cer
```

The **openssl** utility generates a common name **Platform key** for the certificate by setting output format to Distinguished Encoding Rules (DER). DER is a standardized binary format for data encoding.

- Generate a new RSA private key **KEK.key** and a self-signed X.509 certificate **KEK.cer** for the Key Exchange Key database:

```
$ openssl req -quiet \
-newkey rsa:4096 \
-nodes -keyout KEK.key \
-new -x509 -sha256 \
-days 3650 \
-subj "/CN=Key Exchange Key/" \
-outform DER \
-out KEK.cer
```

- Generate a custom certificate **custom\_db.cer**:

```
$ openssl req -quiet \
-newkey rsa:4096 \
-nodes -keyout custom_db.key \
-new -x509 -sha256 \
-days 3650 \
-subj "/CN=Signature Database key/" \
-outform DER \
-out custom_db.cer
```

- Download the Microsoft certificate:

```
$ wget https://go.microsoft.com/fwlink/?LinkId=321194 --user-agent="Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/51.0.2704.103 Safari/537.36" -O MicCorUEFCA2011_2011-06-27.crt
```

- Download the updated forbidden signatures (**dbx**) UEFI Revocation List File for 64 bit system:

```
$ wget https://uefi.org/sites/default/files/resources/x64_DBXUpdate.bin
```

- Use the **google-cloud-cli** utility to create and register the image from a disk snapshot with the required Secure Boot variables:

```
$ gcloud compute images create <example-rhel-10-efi-image> \
--source-image projects/<example_project_id>/global/images/<example_image_name> \
--platform-key-file=PK.cer \
--key-exchange-key-file=KEK.cer \
--signature-database-file=custom_db.cer,MicCorUEFCA2011_2011-06-27.crt \
--forbidden-database-file x64_DBXUpdate.bin \
--guest-os-features="UEFI_COMPATIBLE"
```

- Launch the instance of an **example-rhel-10-efi-image** image with the **Turn on Security Boot** feature from the Google Cloud console.

# Security

Shielded VM and SSH keys

## Shielded VM ?

Turn on all settings for the most secure configuration.

- Turn on Secure Boot ?
- Turn on vTPM ?
- Turn on Integrity Monitoring ?

## Verification

1. Verify if Secure Boot is enabled:

```
$ mokutil --sb-state  
SecureBoot enabled
```

2. Verify the kernel keyring for the custom certificate:

```
$ sudo keyctl list %:platform  
...  
757453569: ---lswrv 0 0 asymmetric: Signature Database key:  
f064979641c24e1b935e402bdbc3d5c4672a1acc  
...
```

## 5.4. CONFIGURING A RHEL INSTANCE ON THE GOOGLE CLOUD MARKETPLACE WITH SECURE BOOT

To ensure that your Red Hat Enterprise Linux (RHEL) instance on Google Cloud has a secured operating system booting process, use Secure Boot. In the Google Cloud Marketplace, you can configure Secure Boot on a RHEL instance launched with the **Turn on Secure Boot** option enabled.

Secure Boot requires Unified Extensible Firmware Interface (UEFI) to verify boot components and protect your system from unauthorized software. The **Turn on Secure Boot** option enables support for the Unified Extensible Firmware Interface (UEFI) boot loader required for Secure Boot. Without UEFI, the Secure Boot feature does not work. By default, it has the Allowed Signature database (**db**) with Microsoft certificates.

### Prerequisites

- You have an account on Google Cloud. For details, see [Google Cloud](#).

### Procedure

1. Launch a publicly available Red Hat Enterprise Linux instance from Google Cloud console.

2. Install the **keyutils** package:

```
$ dnf install keyutils
```

3. Enable the **Turn on Secure Boot** option on the RHEL Google Cloud instance:

## Security

Shielded VM and SSH keys

### Shielded VM ?

Turn on all settings for the most secure configuration.

- Turn on Secure Boot ?
- Turn on vTPM ?
- Turn on Integrity Monitoring ?

## Verification

1. Verify if Secure Boot is enabled:

```
$ mokutil --sb-state
```

```
SecureBoot enabled
```

2. Verify the kernel keyring for the custom certificate:

```
$ sudo keyctl list %:platform
```

```
4 keys in keyring:
12702216: ---lswrv 0 0 asymmetric: Microsoft Corporation UEFI CA 2011:
13adbf4309bd82709c8cd54f316ed522988a1bd4
449796228: ---lswrv 0 0 asymmetric: Red Hat Secure Boot CA 8:
e1c6c580aa1e21d585aad9bf20f3929e5ec1f08b
50338534: ---lswrv 0 0 asymmetric: Red Hat Secure Boot CA 5:
cc6fa5e72868ba494e939bbd680b9144769a9f8f
681047026: ---lswrv 0 0 asymmetric: Microsoft Windows Production PCA 2011:
a92902398e16c49778cd90f99e4f9ae17c55af53
```

# CHAPTER 6. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH INTEL TDX

Intel Trust Domain Extensions (TDX) is a security type of Confidential Virtual Machine (CVM), which provides a secure and isolated environment for VM. This approach is an advancement to the former technology, Intel Software Guard Extensions (SGX).

SGX provides VM isolation from the hypervisor and cloud service providers by creating secure memory regions known as enclaves. Application code stored in enclaves has access to memory and data stored inside the enclave, making it inaccessible to outside entities.

TDX creates hardware-isolated VMs called Trusted Domains (TDs). It ensures that only a VM accesses its memory and TD VMs are isolated from Virtual Machine Manager (VMM), hypervisors, other VMs, and the host. This ensures that while using resources from hypervisor, CPU, TD VMs remain secure by maintaining data confidentiality and integrity.

The main difference between SGX and TDX is that SGX works at application level while TDX works at virtualization level by limiting hypervisor access.



## NOTE

Before deploying Red Hat Enterprise Linux (RHEL) on a public cloud platform, always check with the corresponding cloud service provider for the support status and certification of the particular RHEL instance type.

## 6.1. UNDERSTANDING INTEL TDX SECURE BOOT PROCESS

1. **Initialization and measurement:** A TDX-enabled hypervisor sets the initial state of a VM. This hypervisor loads the firmware binary file into the VM memory and sets the initial register state. The Intel processor measures the initial state of the VM and provides details to verify the initial state of the VM.
2. **Firmware:** The VM initiates the UEFI firmware. The firmware might include stateful or stateless Virtual Trusted Platform Module (vTPM) implementation. Stateful vTPM maintains persistent cryptographic state across VM reboots and migrations, whereas stateless vTPM generates fresh cryptographic state for each VM session without persistence. Virtual Machine Privilege Levels (VMPL) technology isolates vTPM from the guest. VMPL offers hardware-enforced privilege isolation between different VM components and the hypervisor.
3. **vTPM:** Depending on your cloud service provider, for stateful vTPM implementation, the UEFI firmware might perform a remote attestation to decrypt the persistent state of vTPM. The vTPM also gathers data about the boot process, such as Secure Boot state, certificates used for signing boot artifacts, or UEFI binary hashes.
4. **Shim :** When the UEFI firmware finishes the initialization process, it searches for the extended firmware interface (EFI) system partition. Then, the UEFI firmware verifies and executes the first stage boot loader from there. For RHEL, this is **shim**. The **shim** program allows non-Microsoft operating systems to load the second stage boot loader from the EFI system partition.
  - a. **shim** uses a Red Hat certificate to verify the second stage boot loader ( **grub** ) or Red Hat Unified Kernel Image (UKI).

- b. **grub** or **UKI** unpacks, verifies, and executes Linux kernel and initramfs, and the kernel command line. This process ensures that the Linux kernel is loaded in a trusted and secured environment.
5. **Initramfs:** In initramfs, vTPM information automatically unlocks the encrypted root partition in case of full disk encryption technology.
    - a. When the root volume becomes available, **initramfs** transfers the execution flow there.
  6. **Attestation:** The VM tenant gets access to the system and can perform a remote attestation to ensure that the accessed VM is an untampered Confidential Virtual Machine (CVM). Attestation is performed based on information from the Intel processor and vTPM. This process confirms the authenticity and reliability of the initial CPU and memory state of the RHEL instance and Intel processor.
  7. **TEE:** This process creates a Trusted Execution Environment (TEE) to ensure that booting of the VM is in a trusted and secured environment.

## 6.2. CONFIGURING A RHEL INSTANCE ON GOOGLE CLOUD WITH INTEL TDX

Intel TDX is a hardware based trusted execution environment (TEE). This confidential computing technology provides isolation and integrity to virtual machines.

### Prerequisites

- You have installed the **openssh** and **openssh-clients** packages.
- You have installed the **google-cloud-cli** utility. For instructions, see [Installing gcloud CLI on RHEL](#).
- You have launched the instance only from a supported Google Cloud instance type. For details, see [Supported Google Cloud instance types for TDX](#) .

### Procedure

1. Log in to your Google Cloud account by using the **google-cloud-cli** utility:

```
$ gcloud auth login
```

2. Create a new project:

```
$ gcloud projects create <example_tdx_project> --name="RHEL TDX Project"
```

3. Configure the project by setting the **google-cloud-cli** utility property:

```
$ gcloud config set project <example_tdx_project>
```

4. Create a RHEL compute instance:

```
$ gcloud compute instances create <example-rhel-9-tdx-instance> \
--confidential-compute-type=TDX \
--machine-type=c3-standard-4 \
--min-cpu-platform="Intel Sapphire Rapids" \
```

```
--maintenance-policy="TERMINATE" \
--image=<rhel-guest-image-9-6-20250410-0-x86-64> \
--image-project="rhel-cloud" \
--subnet=<example_subnet>
```

5. Open ports on the RHEL instance:

```
$ gcloud compute firewall-rules create allow-ssh \
--allow tcp:22 \
--source-ranges 0.0.0.0/0 \
--target-tags ssh
```

6. Connect to the RHEL Google Cloud instance.

- a. Optional: Create a new key pair:

```
$ gcloud compute ssh <cloud-user>@<example-rhel-9-tdx-instance>
```

- b. Connect to the RHEL instance by using your key pair:

```
$ ssh -i <private key> <cloud-user>@<instance ip>
```

## Verification

- Check the VM configuration of the RHEL instance:

```
$ gcloud compute instances describe <example-rhel-9-tdx-instance> --format="get(metadata)"
```

- Check kernel logs to verify the status of the TDX feature:

```
$ sudo dmesg | grep -i tdx
```

```
[ 0.000000] tdx: Guest detected
[ 1.334504] process: using TDX aware idle routine
[ 1.413419] Memory Encryption Features active: Intel TDX
[ 3.606910] systemd[1]: Detected confidential virtualization tdx.
[ 10.680475] systemd[1]: Detected confidential virtualization tdx.
```

- Check the CPU flags of the VM:

```
$ grep -E 'tdx_guest' /proc/cpuinfo
```

```
flags : ...`tdx_guest`...
```

## Additional resources

- [General purpose machine family for Compute Engine](#)

# CHAPTER 7. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH AMD SEV SNP

AMD Secure Encrypted Virtualization with Secure Nested Paging (SEV-SNP) aims to prevent VM integrity-based attacks and reduce memory integrity violations. AMD processors offer three hardware-based security mechanisms: SEV, SEV-ES, and SEV-SNP.

- **SEV:** The SEV mechanism encrypts virtual machine (VM) memory to prevent the hypervisor from accessing VM data.
- **SEV-ES:** SEV with Encrypted State (SEV-ES) extends SEV by encrypting CPU register states. This mechanism prevents the hypervisor from accessing or modifying VM CPU registers. Despite providing isolation between hypervisor and VM, it is still vulnerable to memory integrity attacks.
- **SEV-SNP:** SEV-SNP is an enhancement to SEV-ES that adds memory integrity protection along with VM encryption. This mechanism prevents the hypervisor from modifying page tables to redirect VM memory access, protecting against replay attacks and memory tampering.



## NOTE

Before deploying Red Hat Enterprise Linux (RHEL) on a public cloud platform, always check with the corresponding cloud service provider for the support status and certification of the particular RHEL instance type.

## 7.1. PROPERTIES OF SEV-SNP

- **Secure Processor:** The AMD **EPYC** processor integrates a Secure Processor (SP) subsystem. AMD SP is a dedicated hardware component to manage keys and encryption operations.
- **Memory Integrity:** For managing virtualization and isolation, memory management unit (MMU) utilizes page tables to translate virtual addresses to guest-physical addresses. SEV-SNP uses nested page tables for translating guest-physical addresses to host-physical addresses. Once nested page tables are defined, the hypervisor or host cannot alter page tables to modify the VM into accessing different pages, resulting in protection of memory integrity. SEV-SNP uses this method to offer protection against replay attacks and malicious modifications to VM memory.
- **Memory Encryption:** The AMD **EPYC** processor hides the memory encryption key, which remains hidden from both host and VM.
- **Attestation report for verification:** A CPU-generated report about RHEL instance information in an authorized cryptographic format. This process confirms the authenticity and reliability of the initial CPU and memory state of the RHEL instance and AMD processor.



## NOTE

Even if a hypervisor creates the primary memory and CPU register state of the VM, they remain hidden and inaccessible to the hypervisor after initialization of that VM.

## 7.2. UNDERSTANDING AMD SEV SNP SECURE BOOT PROCESS

1. **Initialization and measurement:** A SEV-SNP enabled hypervisor sets the initial state of a VM.

This hypervisor loads firmware binary into the VM memory and sets the initial register state. AMD Secure Processor (SP) measures the initial state of the VM and provides details to verify the initial state of the VM.

2. **Firmware:** The VM initiates the UEFI firmware. The firmware might include either stateful or stateless Virtual Trusted Platform Module (vTPM) implementation. Stateful vTPM maintains persistent cryptographic state across VM reboots and migrations, whereas stateless vTPM generates fresh cryptographic state for each VM session without persistence. Virtual Machine Privilege Levels (VMPL) technology isolates vTPM from the guest. VMPL offers hardware-enforced privilege isolation between different VM components and the hypervisor.
3. **vTPM:** Depending on your cloud service provider, for stateful vTPM implementation, the UEFI firmware might perform a remote attestation to decrypt the persistent state of vTPM.
  - a. The vTPM also measures facts about the boot process such as Secure Boot state, certificates used for signing boot artifacts, UEFI binary hashes, and so on.
4. **Shim:** When the UEFI firmware finishes the initialization process, it searches for the extended firmware interface (EFI) system partition. Then, the UEFI firmware verifies and executes the first stage boot loader from there. For RHEL, this is **shim**. The **shim** program allows non-Microsoft operating systems to load the second stage boot loader from the EFI system partition.
  - a. **shim** uses a Red Hat certificate to verify the second stage boot loader (**grub**) or Red Hat Unified Kernel Image (UKI).
  - b. **grub** or **UKI** unpacks, verifies, and executes Linux kernel and initial RAM filesystem (**initramfs**), and the kernel command line. This process ensures that the Linux kernel is loaded in a trusted and secured environment.
5. **Initramfs:** In **initramfs**, vTPM information automatically unlocks the encrypted root partition in case of full disk encryption technology.
  - a. When the root volume becomes available, **initramfs** transfers the execution flow to the root volume.
6. **Attestation:** The VM tenant gets access to the system and can perform a remote attestation to ensure that the accessed VM is an untampered Confidential Virtual Machine (CVM). Attestation is performed based on information from AMD SP and vTPM. This process confirms the authenticity and reliability of the initial CPU and memory state of the RHEL instance and AMD processor.
7. **TEE:** This process creates a Trusted Execution Environment (TEE) to ensure that booting of the VM is in a trusted and secured environment.

## 7.3. CONFIGURING A RHEL INSTANCE ON GOOGLE CLOUD WITH AMD SEV SNP

To create a trusted boot environment, you can configure AMD Secure Encrypted Virtualization with Secure Nested Paging (SEV-SNP) on Red Hat Enterprise Linux (RHEL) instances on Google Cloud. SEV-SNP aims at protecting your data from access by the hypervisor and cloud service provider.

### Prerequisites

- You have installed the **openssh** and **openssh-clients** packages.

- You have installed the **google-cloud-cli** utility. For details, see [installing gcloud CLI on RHEL](#).
- You have created a Google Cloud instance using an AMD EPYC processor-based machine type from the supported list. For details, see [the specified Google Cloud instance types](#).

## Procedure

1. Log in to Google Cloud by using the **google-cloud-cli** utility:

```
$ gcloud auth login
```

2. Create a new Google Cloud project:

```
$ gcloud projects create <example_sev_snp_project> --name="RHEL SEV SNP Project"
```

3. Configure the Google Cloud project:

```
$ gcloud config set project <example_sev_snp_project>
```

4. Create a RHEL compute instance:

```
$ gcloud compute instances create <example-rhel-10-sev-snp-instance> \
--confidential-compute-type=SEVSNP \
--machine-type=n2d-standard-2 \
--min-cpu-platform="AMD Milan" \
--maintenance-policy="TERMINATE" \
--image=<rhel-guest-image-10-0-20251016-6-x86-64> \
--image-project="rhel-cloud" \
--subnet=<example_subnet>
```

5. Connect to the RHEL instance by using public and private RSA key pair:

- a. Connect to the RHEL instance by using a new public and private RSA key pair:

```
$ gcloud compute ssh <cloud_user>@<example-rhel-10-sev-snp-instance>
```

- b. Connect to the RHEL instance by using an existing key pair:

```
$ ssh -i <example_private_key> <cloud_user>@<instance_ip>
```

## Verification

- Check metadata of VM configuration on RHEL instance:

```
$ gcloud compute instances describe <example-rhel-10-sev-snp-instance> \
--format="get(metadata)"
```

- Check kernel logs to verify status of SEV-SNP:

```
$ sudo dmesg | grep -i sev
```

```
[ 0.302688] Memory Encryption Features active: AMD SEV SEV-ES SEV-SNP
```

```
[ 0.303645] SEV: Status: SEV SEV-ES SEV-SNP
[ 0.626662] SEV: APIC: wakeup_secondary_cpu() replaced with wakeup_cpu_via_vmgexit()
[ 0.686825] SEV: Using SNP CPUID table, 57 entries present.
[ 0.687645] SEV: SNP running at VMPL0.
[ 1.477045] SEV: SNP guest platform devices initialized.
[ 19.921893] systemd[1]: Detected confidential virtualization sev-snp.
[ 25.604801] systemd[1]: Detected confidential virtualization sev-snp.
[ 29.089032] sev-guest sev-guest: Initialized SEV guest driver (using VMPCCK0
communication key)
```

## Additional resources

- [General purpose machine family for Compute Engine](#)

# CHAPTER 8. CONFIGURING RHEL ON PUBLIC CLOUD PLATFORMS WITH UKI

To ensure that a Red Hat Enterprise Linux (RHEL) instance has a secured boot process from an untrusted storage such as confidential virtual machine (CVM) on a public cloud platform, use Unified Kernel Image (UKI).

## 8.1. INTRODUCTION TO UNIFIED KERNEL IMAGE

To extend the secure boot protection throughout the entire boot chain, use Unified Kernel Image (UKI).

### Components of UKI

Unified Kernel Image (UKI) is a Unified Extensible Firmware Interface (UEFI) Portable Executable (PE) binary for the UEFI environment that bundles together the essential components of an operating system. UKI binary components extend the Secure Boot coverage by including the **initramfs** and the kernel command line. **Initramfs** is a part of the Linux startup process, while the kernel command line gives you limited access to define parameters. The primary components contained within the UKI binary are:

- The **.linux** section stores the Linux kernel image.
- The **.initrd** section stores the initial RAM filesystem **initramfs**.
- The **.cmdline** section stores the kernel command line.
- Additional sections, such as **.sbat**.
- The Red Hat signature.

### Features of RHEL UKI with pre-builtinitramfs

- Prohibits any malicious agent or component to alter any objects in the boot chain.
- Due to pre-built **initramfs**, the user does not need to build its custom **initramfs**, which results in a faster kernel installation.
- Provides support for the pre-built **initramfs** systems as it is similar in all installations such as virtual machine (VMs), containers, or cloud instances.
- Provides support for the **x86\_64** architecture.
- Includes the **kernel-uki-virt** package.
- Built for virtual machines and cloud instances.

### Limitation of UKI because of the reduced flexibility of the boot process

- When building the UKI, the operating system vendor creates **initramfs**. As a consequence, the listed and included kernel modules are static. You can use the **systemd** system and configuration extensions to address this limitation.
- The kernel command line parameters are static, which limits the use of parameters for different instance sizes or debugging options.

You can use the UKI command line extensions to overcome this limitation.

## 8.2. UNDERSTANDING THE UKI SECURE BOOT PROCESS

To protect your system against unauthorized boot-time modifications, use the secure boot mechanism with Unified Kernel Image (UKI).

When using UKI with secure boot, the system verifies each component in the boot chain to ensure system integrity and prevent malicious code execution.

### Procedure

1. **UEFI Firmware:** The boot process starts from the Unified Extensible Firmware Interface (UEFI) firmware. For boot, Red Hat Enterprise Linux (RHEL) UKI requires UEFI firmware, because legacy basic input/output system (BIOS) firmware is not supported.
2. **Shim boot loader:** Use the **shim** boot loader for booting rather than directly booting the UKI (a PE binary) from the UEFI firmware. **shim** includes additional security mechanisms such as Machine Owner Key (**MOK**) and Secure Boot Advanced Targeting (SBAT).
3. **Signature verification (Secure Boot UEFI mechanism)** During boot, **shim** reads the UKI binary and the secure boot UEFI mechanism verifies the signature of UKI against trusted keys stored in the Secure Boot Allowed Signature Database (**db**) of the system, the **MOK** database, and the built-in database of the **shim** binary. If the signature key is valid, the verification passes.
4. **SBAT verification:** Immediately after signature verification, the **shim** boot loader verifies the SBAT rules at startup.

During SBAT verification, the system compares generation numbers, for components such as **systemd.rhel** or **linux.rhel**, embedded in the UKI by using the **.sbat** section against values in the **shim** boot loader. If the generation number for a component in the **shim** is higher than the generation number in the UKI, the binary gets automatically discarded, even if it was signed by a trusted key.

Note that the generation number is a version identifier for UEFI applications, such as **shim** and **grub**.

5. **Unpacking and Execution:** If verification passes, control passes from **shim** to the **systemd-stub** code in the UKI to continue the boot process.
6. **systemd-stub add-ons:** During execution, **systemd-stub** unpacks and extracts the contents of the **.cmdline** section, the plain text kernel command line, and the **.initrd** section, the temporary root file system, for the boot process.  
Note that **systemd-stub** reads UKI add-ons, which are also PE binaries, and verifies their signatures to safely extend the kernel command line of UKI by appending the **.cmdline** content from add-ons. **systemd-stub** reads add-ons from two locations:
  - Global (UKI-independent) add-ons from the **/loader/addons** directory on the Extensible Firmware Interface (EFI) System Partition (ESP).
  - Per-UKI add-ons from the **/EFI/Linux/<UKI-name>.extra.d** directory on the ESP.
7. Control passes from **systemd-stub** to the Linux kernel and the operating system boot process continues.  
From this point, secure boot with UKI mechanisms follows the standard kernel boot process.

